



Systems programming

6 – pipes



MEEC LEEC MEAer LEAer MEIC-A

João Nuno Silva



TÉCNICO
LISBOA

Bibliography

- Pipes - Summary, M. D. McIlroy
- The Linux Programming Interface, Michael Kerrisk
 - Chapter 44
- Beej's Guide to Interprocess Communication, Brian Hall
 - Chapters 4-5
- The Linux Programmer's Guide - Sven Goldt, Sven van der Meer, Scott Burke
 - Sections 6.1 .. 6.3

- M. D. McIlroy - October 11, 1964
 - We should have some ways of coupling programs like garden hose
 - screw in another segment when it becomes necessary to message data in another way
 - This is the way of IO also.
- In Unix pipes are the original inter-process communication mechanisms
 - They are anonymous and allow related processes to exchange data
 - It is also deployed by the shell
 - Uses simple communication protocol (just text)

Pipes in the shell


- pipe operation
 - connects the standard output of one program to the standard input of another.
 - A chain of programs connected in this way is called a pipeline.
- `ls | wc`
 - Counts character/word/line count for the current directory listing
- `tr -c '[:alnum:]' '\n*' | sort -iu |
grep -v '^[0-9]*$'`

Pipes in the shell

- All the stages in a pipeline run concurrently.
 - Each stage waits for input on the output of the previous one,
 - no stage has to exit before the next can run.
- It is unidirectional.
 - $p1 \mid p2$ $p1 \rightarrow p2$ $p1 \leftarrow p2$
 - Impossible to pass information back
 - Just $p2$ exit notification

Pipes

- Read/Writes
 - File operations
- Processes
 - Should be related
 - Father/soon
 - Brothers



Pipes

- Pipe creation:

```
int pipe(int fd[2]);  
int pipe2(int pipefd[2],  
          int flags);
```
- Opens two files
 - **fd[0]** descriptor open for reading
 - **fd[1]** descriptor open for writing
- Returns
 - 0 successful
 - 1 unsuccessful
- Pipes can only
 - connect processes with a common ancestor
- Pipe
 - is managed as a open file

Pipes

- Communication (data read/write)
 - `ssize_t read(int fd, void *buf, size_t count);`
 - `ssize_t write(int fd, void *buf, size_t count);`
- 1st argument
 - File descriptor (fd[0] or fd[1])
- 2nd argument
 - data buffer address (data destination/source)
- 3rd argument
 - number of bytes to read/write
- Return number of bytes read/written



Closing Pipes

- Closing all write ends
 - Read will return 0
- Closing all read ends
 - Write will produce an exception
- Processes should close not needed ends
 - For previous notifications to work

Pipes

- Pipes offer some primitive synchronization
 - Process wait for the others
- If a process attempts to read from an empty pipe,
 - then **read** will block until data is available.
- If a process attempts to write to a full pipe,
 - then **write** blocks until sufficient data has been read from the pipe to allow the write to complete.

Pipes

- The communication channel provided is
 - byte stream:
- there is no concept of message boundaries.
- Reads returns all available data up to the size of buffer
 - **write(10 bytes) → write(10 bytes) → read (13 bytes)**
 - Is possible, leaves 7 bytes in the pipe

Pipes

- Nonblocking I/O is possible
 - using `O_NONBLOCK` status flag
 - reads/writes return immediately
 - Programmer must verify if data was transferred
- Writes on pipe are atomic:
 - the output data is written to the pipe as a contiguous sequence.
 - Even if two writes are concurrent
 - In different processes

Pipes

- A pipe has a limited capacity.
 - Data is stored in the operating system
 - Until read, pip destroyed or reboot
 - Different OS have different limits for the pipe capacity
 - POSIX.1-2001 requires PIPE_BUF to be at least 512 bytes.
 - On Linux PIPE_BUF is 4096 bytes.
 - Applications should not rely on a particular capacity
 - application should consume data as soon as possible

Pipes

- Messages are limited to byte streams.
- Information flow is unidirectional
 - One process reads one process writes
 - Uses file descriptors functionality
- Major limitation
 - Processes should be related
- How to implement pipes that are accessible by other processes?
 - Giving them a name
 - Registering them in the File system



Named pipes / FIFOs

FIFO / Named Pipes

- To solve Pipes limitations
 - FIFOs were defined
 - Also referred as named pipes
- Can be used by unrelated processes
- Are referred and identified by a file in the file system
- A FIFO is special file similar to a pipe,
 - That is created in a different way
 - Instead of being an anonymous communications channel,
 - FIFO is entered into the file system by calling `mkfifo()`
 -

FIFO / Named Pipes

- FIFO creations
 - `int mkfifo(const char *pathname, mode_t mode);`
 - 1st argument - FIFO name (full path)
 - 2nd argument - Access permissions (like a regular file)
- On success **mkfifo()** returns 0.
 - In the case of an error, -1 is returned (errno is set appropriately).
- Once a FIFO is created a special file is available
 - any process can open it for reading or writing, in the same way as an ordinary file.

FIFO / Named Pipes

- Before being used the FIFO should be opened

int open(const char *pathname, int flags);

- Like a regular file
- 1st argument - FIFO name
- 2nd argument - Bits that define access mode
 - O_RDONLY (just reading)
 - O_WRONLY (just writing)
 - O_NONBLOCK (non blocking I/O)
- The return value is
 - -1 in case of error
 - Or a positive file descriptor

FIFO / Named Pipes

- A FIFO has to be opened at both ends simultaneously
 - before any process returns from open
 - Opening a FIFO for reading normally blocks until some other process opens the same FIFO for writing, and vice versa.
- After opening, FIFO offer the same synchronization as pipes
 - empty/full FIFO
- Data access is similar to a pipe
 - Data stream and atomicity

Next on PSIS

- sockets