



Systems programming

7 – Sockets




MEEC LEEC MEAer LEAer MEIC-A

João Nuno Silva






Bibliography

- Beej's Guide to Unix IPC
 - Chapter 11
 - The Linux programming interface
 - Chapter 56, 57, 57
 - Solaris SUN Programming Interfaces Guide
 - Ch 6, 7
- 



Pipes / Names Pipes

- Only local Communication
 - Impossible “private” communication
 - Impossible one to one
 - Impossible directed
 - Only one protocol implemented
 - Stream communication
 - No messages boundaries
- 

Sockets


- Transparency
 - Communication inter/intra machines is the same
- Compatibility
 - With existing communication mechanism
-
- Stream oriented
- Message Oriented
- Use of File system mechanisms

Sockets

- Introduced in 1981 on BSD 4.1
- It is an API
 - that define access points to applications
 - following the client-server architecture
- Sockets programming
 - more complex than files
 - More parameters
 - More system calls
- Main difference between FS base and socket based communication
 - How channels are opened and created.



Sockets

- Transparency
 - Compatibility
 - Use of files
 - to reference channels
 - Use of the Regular I/O API
 - Send/receive data
- 



Sockets types



Message oriented Sockets


- Each write is a message
- No message interleaving in the channel
- Each message is read atomically
- No data interleaving
 - Concurrent writes to not affect each other
- A read is concluded only after the conclusion of the write
- Two writes imply two reads

Message oriented sockets

Definition of a communication point	socket	Telephone
Assignment of a address to a communication point	Bind /address	Assignment of phone number
Send message	sendt()	Send SMS
Receive message	recvfrom(Receive SMS
End of communication	close()	Turn off phone



Stream oriented Sockets

- Multiple writes form a stream/flow of data
 - No data interleaving in the channel
 - Each piece of data is processed atomically
 - No data interleaving
 - Concurrent writes do not affect each other (system serializes them)
 - Reads are serialized
 - 2nd read starts after the first
- 

Stream oriented sockets

Definition of a communication point	socket	Telephone
Assignment of a address to a communication point	Bind /address	Assignment of phone number
Listen to incoming connections	listen	connection of a phone to the network
Start connection	connect()	phone call initiator dials destination number
Receiver established connection	accept()	receiver accepts call lifting the handset
Send / receive of data	send(),recv()	talk
End of communication	close()	lowering of handset

Message Reception

- On other IPC how receives messages?
 - Any process that open the channel
- Does the sender know the identity?
 - No
- How to solve
 - Assign each channel an address
 - Only one process can read “from” one address

Socket Types

- The socket type determines
 - The characteristics of communication
 - Delivery guarantees, ordering guarantees, communication directions
- Are defined as constants starting with the SOCK_ prefix
- **SOCK_STREAM**
 - stream oriented socket / connection oriented
- **SOCK_DGRAM**
 - Message oriented sockets / datagram socket / connectionless

Common socket types

- SOCK_DGRAM
 - Non reliable delivery
 - Packets can be lost or changed
 - No order guarantee
 - Non existing connection
 - Application should define recipient address for each message
 - (Uni/Bi)directional
 - Recipient can retrieve sender address and reply
- SOCK_STREAM
 - Reliable delivery
 - Ordered delivery
 - 1st packet to be sent is the 1st to be received
 - Connection oriented
 - Connection setup required before sending messages
 - Bidirectional by default



Socket domains



Socket Domains

- The same API allows the creation of different sockets
- AF_UNIX
 - Communication between processes in the same machine
- AF_INET
 - Communication between processes in different machines
 - IPv4
- AF_INET6
 - Communication between processes in different machines
 - IPv6

Socket Domains

- Determines the nature of the communication (local/LAN/WAN)
- Determine the format of the addresses
 - AF_UNIX – a string
 - AF_INT – 4 bytes

Socket address

- Some operations require an address
 - Bind / connect
 - Sendto / recvfrom
- `struct sockaddr *src_addr`
 - Placeholder for various address classes
 - `sockaddr_un` - unix
 - `sockaddr_in` - IP
 - `sockaddr_nl` - netlink
 - `sockaddr_atalk` - appletalk

```
struct sockaddr {  
    sa_family_t  
    sa_family;  
    char  
    sa_data[14];  
}
```

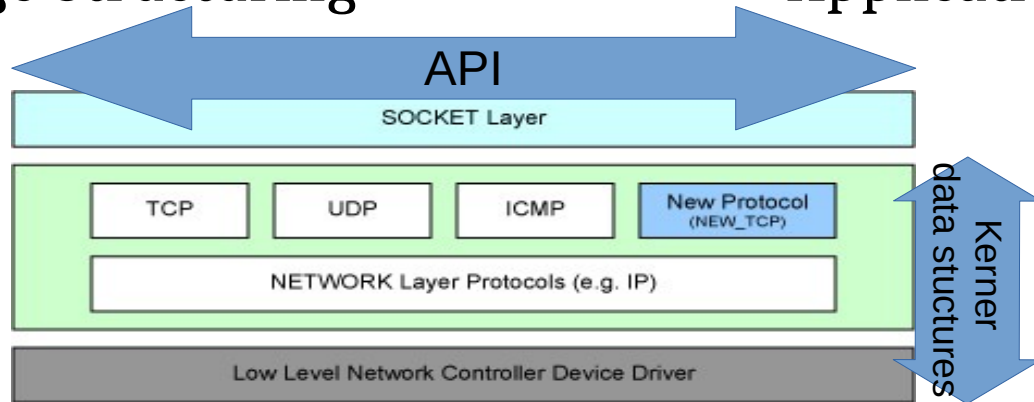
Protocol

- The communication protocol
 - Depends on the socket Type and domain
 - Not all combinations are possible
 - Defines communication characteristics

		Domain		
		AF_UNIX	AF_INET	AF_UNSPEC
Type	SOCK_STREAM	YES	TCP	SPP
	SOCK_DGRAM	YES	UDP	IDP
	SOCK_RAW		IP	Sim
	SOCK_SEQPACKET	YES		SPP

Protocol

- Defines
 - Addressing
 - Delivery guarantees
 - Message Structuring
- Affects
 - Kernel data structures
 - API
 - Application






Sockets functioning




Datagram sockets

- Client

- Socket creation
-  Address assignment
- Message sending
 - Explicit address

- Server

- Socket creation
-  Address assignment
- Reception of message

Datagram sockets

- Client

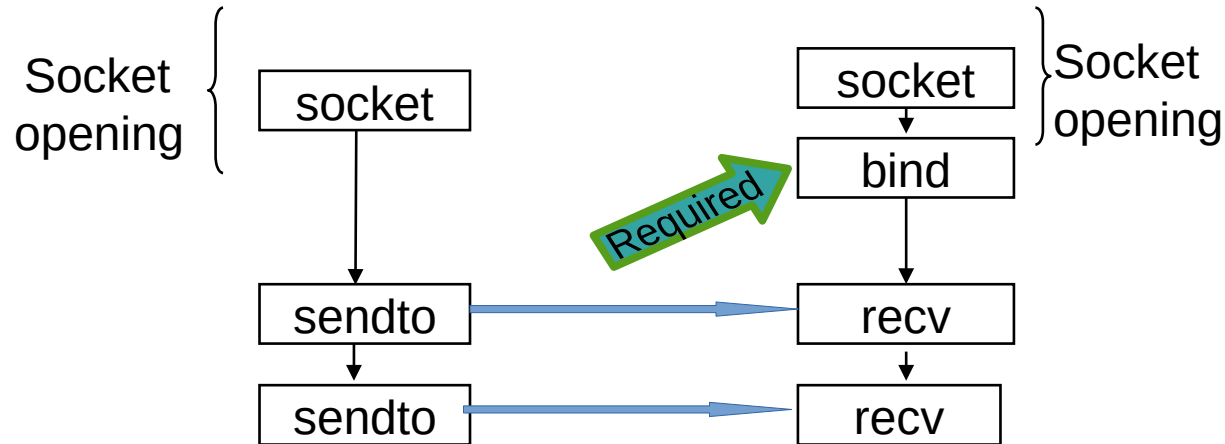
- Socket creation
- Address assignment
- Message sending
 - Explicit address

Optional

- Server


- Socket creation
- Address assignment
- Reception of message

Required




Datagram sockets

- Server

- Socket creation
-  – Address assignment
- Reception of message

- Client

-  Socket creation
- Address assignment
- Message sending
 - Explicit address


Datagram sockets

- Server

- Socket creation
- Address assignment
- Reception of message

- Client

- Socket creation
- Address assignment
- Message sending
 - Explicit address



Datagram sockets

- message oriented
- Without connection
 - Client just sends messages to server
 - Needs to address each message
- Without guarantees
 - Delivery
 - Order
 - integrity
- On the client
 - Create socket
 - Send messages
- On the server
 - Create socket
 - Assign address
 - Receive messages
 - One socket can receive messages from multiple clients
- Message transmission
 - sendto
 - Recv / recvfrom



Stream oriented sockets

Client

- Socket creation
- Connection to server
 - Explicit address
- Message sending
- Close

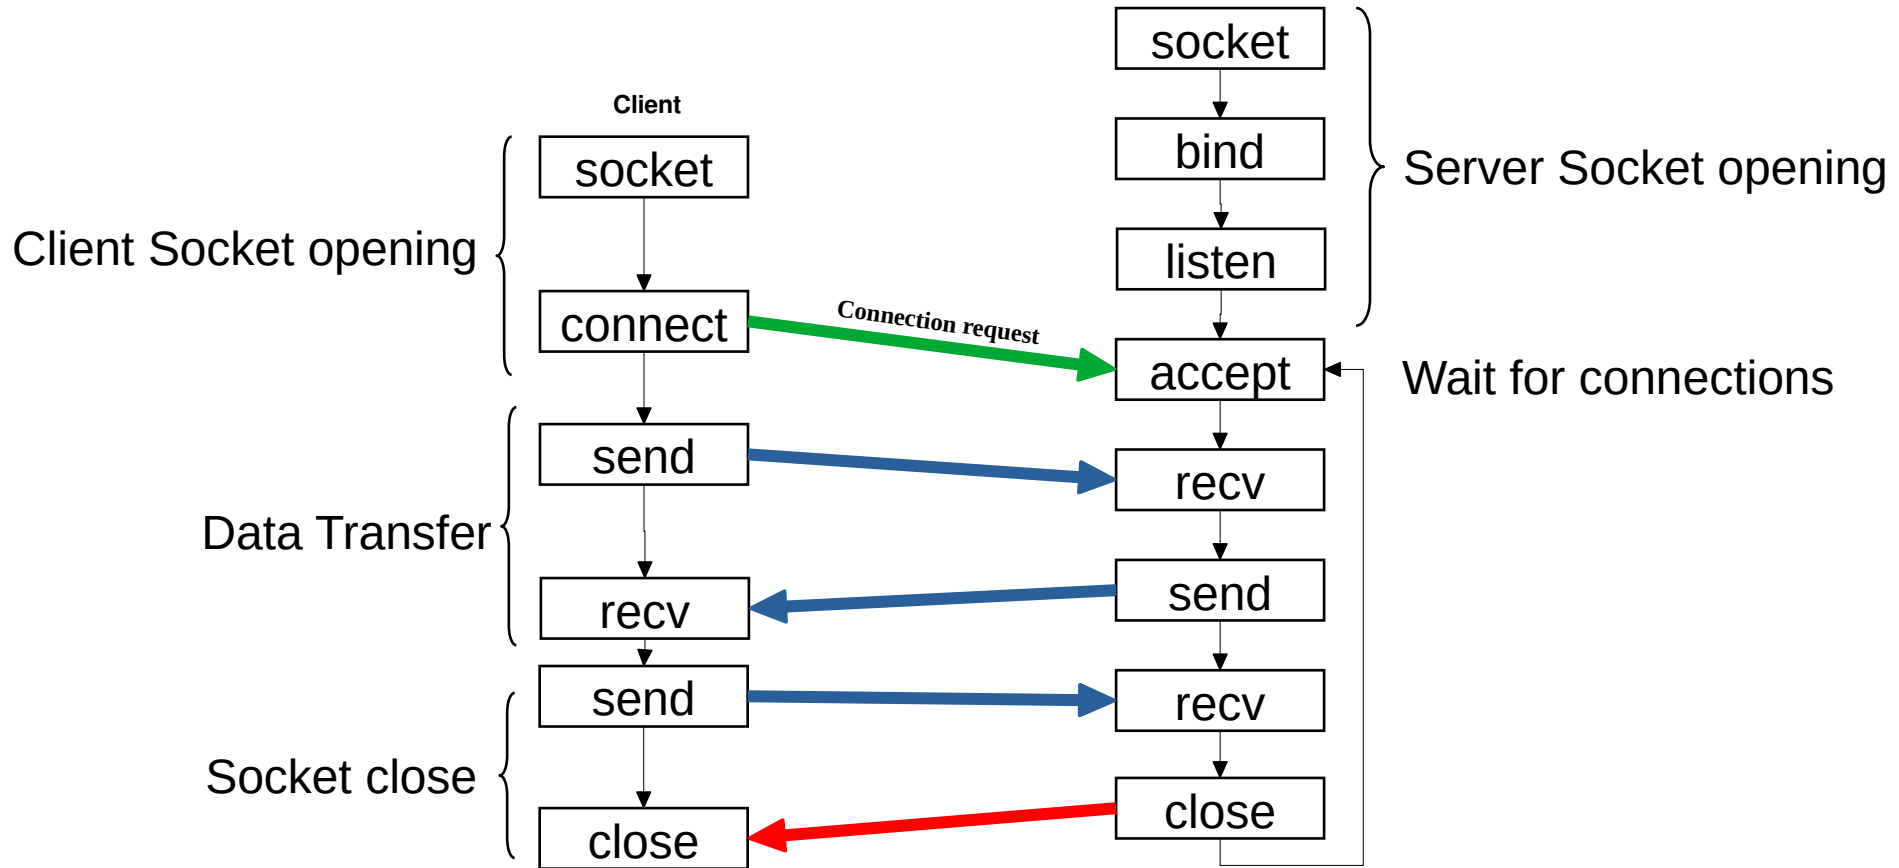
Server

- Socket creation
- Address assignment
- Reception of connections
 - Connection acceptance
- Reception of message
- Close

Stream oriented sockets

- With connection
 - Client should connect to the server
 - Does not need to address each message
- With guarantees
 - Delivery
 - Order
 - integrity
- Connection establishing
 - Listen (server)
 - Connect (Client)
 - Accept (Serve)
- On the server
 - A new socket is created
 - Dedicated to communication with the client
 - Original socket can receive more connections
- Message transition
 - read/recv
 - Write/send
 - Sendto/recvfrom
 - Not needed

Stream oriented sockets



Unix vs Internet sockets

UNIX

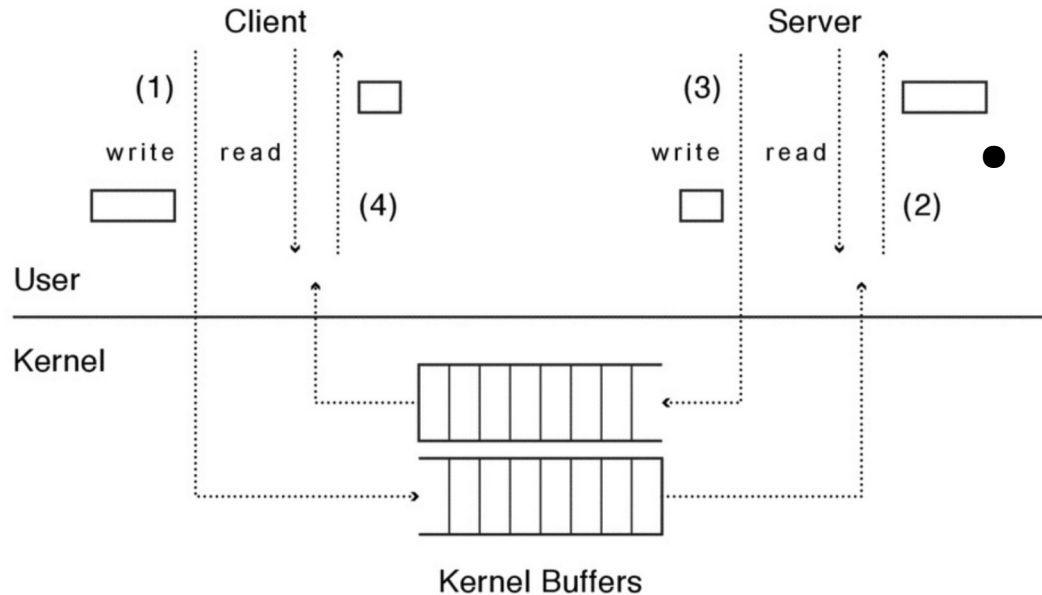
- Addresses
 - String (file name)
- Data transmission
 - App → kernel → App
- Fault tolerance
 - All data delivered

Internet

- Addresses
 - IP address + port
- Data transmission
 - App → kernel → network → kernel → App
- Fault tolerance
 - Depends on the protocol
- Data heterogeneity


Unix sockets data transmission

- Data is sent to kernel
 - send/sendto
- Data is stored in the kernel
 - Until other process read it




○ IMP △ PLURIBUS IMP
 □ TIP ~~~ SATELLITE CIRCUIT

NAMES SHOWN ARE IMP NAMES, NOT (NECESSARILY) HOST NAMES





Unix sockets data transmission

- Data is sent to kernel
 - send/sendto
 - Data is stored in the kernel
 - It can be forwarded to other computer
 - Protocols guarantee that messages
 - Transverse the whole Internet
 - Reach the correct computer
- 

Sockets fault tolerance

UNIX

- All data delivered
- Kernel stores all data until it is read

Internet

- Message oriented sockets
 - Messages can be lost
 - Message order is not guaranteed
 - Reduce message processing overhead
- Stream oriented message
 - Add data is transmitted and in the correct order
 - High communication overhead



Sockets addresses



UNIX sockets addresses

- Unix domain addresses are defined using
- **sockaddr_un** data type
- Definition of the domain
- **sun_family = AF_UNIX**
- Definition of the socket path
- **Strcpy(addr.sun_path, "/tmp/sock_1")**

```
#include <sys/un.h>
```

```
struct sockaddd_un addr;
```

```
addr.sun_family = AF_UNI;
```

```
strcpy(addr.sun_path, "/tmp/sock_1");
```

Internet sockets addresses

- Client needs to identify server
 - On a remote machine
 - From multiple services on such machine
- Internet domain addresses are defined using
 - Sockaddr_in data type
- Definition of the domain
 - sun_family = AF_INET


```
#include <netinet/in.h>
```

```
sockaddr_in local_addr;
```

```
i_addr.sin_port = htons( 22);
```

```
i_addr.sin_addr.s_addr = INADDR_ANY;
```

```
inet_aton("146.193.41.1", &i_addr.sin_addr)
```



Internet sockets addresses

- A service is identified by
 - Network address + Port
- The transmission/reception of data is made using a
 - Port - HW and operating System
 - Socket – App;lication
- A socket can be connected one of 64K ports.
- The first 1K ports (1-1023) are reserved to official services
 - specific services (listed in /etc/services)
 - 22 : SSH 53 : DNS 80 : WWW 115 : secure FTP 443 : secure WWW
- http://en.wikipedia.org/wiki/List_of_TCP_and_UDP_port_numbers
- Ports in between 49152-65535 should not be used by servers
 - Are assigned dynamically to client sockets



Review

Datagram vs stream

- Datagram

- 1 socket multiple clients
- Disconnect
 - Impossible
- Messages limited size
- No guarantees
 - ordering
 - Delivery
- lightweight

- Stream

- 1 socket 1 client
- Disconnect
 - acknowledged
- Streams of unlimited size
- Guarantee
 - Ordering
 - delivery
- heavy weight

Common Errors

- Common errors with socket programming:
 - Incorrect byte ordering
 - Not calling `hton()` e `ntoh()`.
 - Disagree on data size and limits (fix or variable).
 - Non initialization of `len` (`recvfrom`, `accept`)
 - Locks
 - Application level protocol not well defined

Blocking calls

- Most socket API calls are blocking
 - When the process calls such function it gets blocked waiting for an event.
 - Accept - Waits for a connection to be received
 - Connect - Waits for the server to accept the connections
 - recv,recvfrom - Waits for data to be received
 - send,sendto - Waits for data to be transmitted to a lower layer.
- In simple application blocking is good:
 - Adds synchronization
 - Limits resource usage (avoids active wait)

Blocking calls

- In complex applications blocking is a problem:
 - Multiple connections are impossible
 - Simultaneous send/receive are difficult
- There are several solutions
 - Multi-programming
 - Several processes or several threads
 - More complex programming
 - Synchronization required
 - Turn off blocking
 - More complex programming
 - Active wait
 - Use select
 - Wait on multiple descriptors
 - Serializes communication

Blocking calls

```
int select(int, fd_set *, fd_set *, fd_set *, struct timeval *);
```

- 1st parameter
 - Identifies the number of the highest descriptor + 1.
- 2nd parameter
 - array of reading descriptors (if set select verifies if such descriptors have information to be read)
- 3rd parameter
 - array of writing descriptors (if set select verifies if such descriptors can be written to).
- 4th parameter
 - array of “exceptions” (if set select verifies if such descriptors has exception).
- 5th parameter defines the waiting interval (NULL to infinite wait).
- The function returns the number of affected descriptors Return -1 in case of error

Blocking calls

- descriptors are referred in a bit array of type **struct fd_set**.
- Auxiliary functions:
 - `void FD_ZERO(fd_set *); /* create array */`
 - `void FD_CLR(int, fd_set *); /* set bit to 0 */`
 - `void FD_SET(int, fd_set *); /* set bit to 1 */`
 - `int FD_ISSET(int, fd_set *); /* return bit value */`
- The select function activates the correct bits on the correct arrays, depending
 - on the input array
 - on the state of the descriptor



Multiple clients

- Select
 - Fork
 - FD are Inherited
 - Read retrieve messages from same socket
 - Accept can be done in multiple processes
 - Threads
- 