

Systems programming

8 – Indirect communication


MEEC LEEC MEAer LEAer MEIC-A

João Nuno Silva





Bibliography

- Distributed Systems Concepts and Design, George Coulouris, Jean Dollimore, Tim Kindberg, Gordon Blair
 - Chapter 6
 - The Linux Programming Interface, Michael Kerrisk
 - Chapter 52.1 .. 52.5
- 



Point-to-point communication

- Pipes, Fifos, sockets
 - Participants need to exist at the same time
 - Participants need to know address of each other and identities
 - It is necessary to establish communication
-
- Not a good way to communicate with several participants
- Not a good way to implement complex communication



Indirect communication

- Space uncoupling
 - No need to know identity of receiver
 - Anonymous senders
 - Participants can be replaced, updated, replicated, or migrated
- Time uncoupling
 - independent lifetimes
 - requires persistence in the communication channel
- Communication through an intermediary
 - No direct coupling between the sender and the receivers



Indirect Communication

- Scenarios where users connect and disconnect very often
 - Mobile environments, messaging services, forums
- Event dissemination where receivers may be unknown and change often
 - RSS, events feeds in financial services
- Scenarios with very large number of participants
 - Google Ads system, Spotify
- Commonly used in cases when change is anticipated
 - need to provide dependable services

Indirect Communication

- Requires a middleware
- Performance overhead introduced by
 - adding a level of indirection
 - Implementation of reliable message delivery
 - ordering grantees
- More difficult to manage due to lack of direct coupling
- Difficult to achieve end-to-end properties
 - Real time behavior
 - Security

Indirect communication

	<i>Time-coupled</i>	<i>Time-uncoupled</i>
<i>Space coupling</i>	<p><i>Properties:</i> Communication directed towards a given receiver or receivers; receiver(s) must exist at that moment in time</p> <p><i>Examples:</i> Message passing, remote invocation (see Chapters 4 and 5)</p>	<p><i>Properties:</i> Communication directed towards a given receiver or receivers; sender(s) and receiver(s) can have independent lifetimes</p> <p><i>Examples:</i> See Exercise 15.3</p>
<i>Space uncoupling</i>	<p><i>Properties:</i> Sender does not need to know the identity of the receiver(s); receiver(s) must exist at that moment in time</p> <p><i>Examples:</i> IP multicast (see Chapter 4)</p>	<p><i>Properties:</i> Sender does not need to know the identity of the receiver(s); sender(s) and receiver(s) can have independent lifetimes</p> <p><i>Examples:</i> Most indirect communication paradigms covered in this chapter</p>




Message Queues

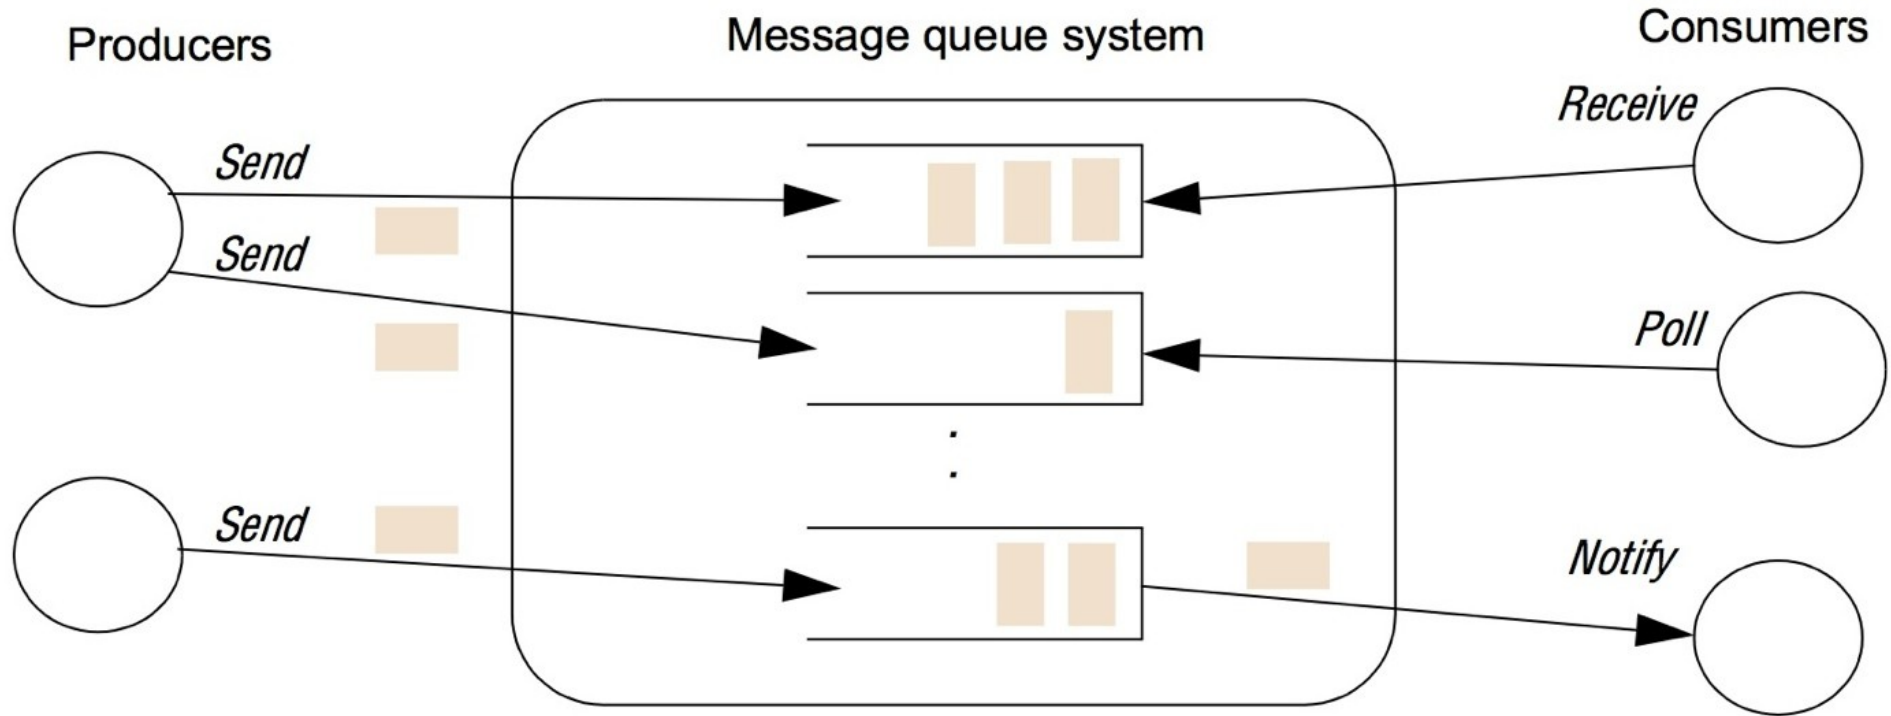




Message Queues

- Intermediary between producers and consumers of data
 - Point-to-Point, not one-to-many
 - Supports time and space uncoupling
 - Programming model
 - producer sends message to specific queue
 - consumers can
 - Block
 - Non-block (polling)
 - Notify
- 

Message Queues



Message Queues

- Many processes can send to a queue
- Many can remove from it
- Queuing policy:
 - usually FIFO
 - Can also be priority-based
- Consumers can select based on metadata



Message Queues

- Messages are persistent
 - Stored until removed (on disk)
- Transaction support:
 - all-or-none operations
 - All message is completely delivered
- Automatic message transformation:
 - on arrival, message transforms data from one format to another (data heterogeneity)

POSIX Message Queues

- Array of bytes
- Priority / message selection
 - API
- Each message has an associated priority,
- Messages are always delivered to the receiving process highest priority first.
- Message priorities range
 - From 0 (low) to `sysconf(_SC_MQ_PRIO_MAX) - 1` (high).
 - On Linux, `sysconf(_SC_MQ_PRIO_MAX)` returns 32768,
 - POSIX.1 requires a range from 0 to 31

POSIX MQ – creation

- mq_open - open a message queue
`mqd_t mq_open(const char *name, int oflag, mode_t mode, struct mq_attr *attr);`
- Name
 - identifier
- Oflags
 - **O_CREAT** | O_RDONLY | O_WRONLY | O_RDWR
- Mode
 - File access modes rwx / ugw 0666

POSIX MQ creation attributes

- 3rd argument of `mq_open`
 - Attributes of the message queues

```
struct mq_attr queue_attr;
```

```
    queue_attr.mq_maxmsg = 16;
```

```
    queue_attr.mq_msgsize = 128;
```

- Posix MQ are implemented as
 - Arrays of messages
 - Circular buffer

POSIX MQ – opening

- mq_open - open a message queue
mqd_t mq_open(const char *name, int oflag)
 - Just 2 arguments
- Default settings
 - Name
 - identifier
 - Oflags -
 - O_RDONLY O_WRONLY O_RDWR

POSIX MQ lifecycle

- **mq_open**
 - Creates / open message queue
- **mq_close**
 - close a message queue descriptor
- **Message queue is assigned to a file**
 - In /dev/msgqueue/
 - File name is used by other processes
 - File object can be accessed
- **mq_unlink**
 - removes a message queue
 - Deletes the file

POSIX MQ - Message Structure

- Array of bytes
- Each message has an associated priority,
- Messages are always delivered to the receiving process highest priority first.
- Message priorities range
 - From 0 (low) to `sysconf(_SC_MQ_PRIO_MAX) - 1` (high).
 - On Linux, `sysconf(_SC_MQ_PRIO_MAX)` returns 32768,
 - POSIX.1 requires a range from 0 to 31

POSIX MQ - write

```
int mq_send(mqd_t mqdes,  
            const char *msg_ptr, size_t msg_len,  
            unsigned int msg_prio);
```

- Writes a message to the queue
- Parameters
 - mqdes
 - queue id (returned from mq_open)
 - Message (byte array) + size
 - msg_priority

POSIX MQ - read

```
ssize_t mq_receive(mqd_t mqdes,  
                  char *msg_ptr, size_t msg_len,  
                  unsigned int *msg_prio);
```

- Reads a message from the queue
 - mqdes – queue id (returned from mq_open)
 - Message + buffer size
 - msg_prio - priority of received message)
- removes the oldest message with the highest priority
 - places it in the buffer pointed to by msg_ptr.

POSIX MQ - read

```
ssize_t mq_receive(mqd_t mqdes,  
                  char *msg_ptr, size_t msg_len,  
                  unsigned int *msg_prio);
```

- If msg_priority not NULL
 - Used to store the priority of received message

Read/write

mq_receive

- If queue is empty
 - Receive blocks until
 - There is is a message
- **mq_timedreceive**
- Blocks until
 - Timeout
 - There is is a message

mq_send

- If queue is full
 - Send blocks until
 - queue has space for message

mq_timedsend

- Blocks until
 - Timeout
 - queue has space for message

POSIX MQ - limits

- Number of messages
- Size of each message
- Values limited by the OS
 - /proc/sys/fs/mqueue/
- Changed on:
 - /etc/security/limits.conf
- On the user program
 - `queue_attr.mq_maxmsg = 16;`
 - `queue_attr.mq_msgsize = 128;`



Group communication



Group communication


- Offers a service where
 - a message is sent to a group
 - then this message is delivered to all members of the group.
- Characteristics
 - Sender is not aware of the identities of the receivers
 - Represents an abstraction over multicast communication
- Added value in terms of
 - Managing group membership
 - Detecting failures and providing reliability and ordering guarantees

Group communication


- Reliable dissemination of information to potentially large numbers of clients,
 - financial industry that require accurate and up-to-date access to data
- Collaborative applications
 - where events must be disseminated to multiple users (e.g. multiuser games)
- Faulty tolerant message delivery
 - Consistent update of replicated data
 - Implementation of highly available (replicated) servers
- System monitoring and management
 - including for example load balancing strategies

Group Communication

- Central abstraction:
 - group & associated membership
- Processes join (explicitly) or leave (explicitly or by failure)
- Send single message to the group of N, not N unicast messages



PSis - 8 - Indirect communication






Process groups

- Abstraction
 - resilient process
- Messages delivered to a process endpoint, no higher
- Messages
 - unstructured byte arrays
 - no marshaling etc
- Level of service
 - socket





Object groups

- Higher level approach
 - Collection of objects (same class!)
 - process same invocations
 - Replication can be transparent to clients
 - Invoke on single object (proxy)
 - Requests sent by group communication
 - Voting in proxy usually
- 

Group membership


- Closed
 - Cooperating servers
 - Internal messages
- Open
 - Notification of services



Group membership

- Interface for group membership changes
 - Operations to create and destroy process groups
 - add or withdraw a process to or from a group
- Failure detection
 - The service monitors the group members
 - Verify if they crash or become unreachable due to communication failure
- Notifying members of group membership changes
 - Notification of group's members when a processes are added, or removed
- Group address expansion
 - Replication of message set to group to all the members

Group membership



Communication Reliability

- Strong delivery reliability properties
 - Delivery integrity - message received same as sent, never delivered twice
 - Delivery validity - outgoing message eventually delivered
- Group communication reliability properties build on Unicast
 - Delivery integrity
 - Deliver message correctly at most once to group members
 - Delivery validity
 - message sent will be eventually delivered (if not all group members fail)
 - Agreement/consensus
 - Delivered to all or none of the group members

Message ordering

- FIFO ordering
 - first-in-first-out from a single sender to the group
- Causal ordering
 - preserves potential causality, happens before
- Total ordering
 - messages delivered in same order to all processes
- Perspective
 - Strong reliability and ordering is expensive: scale limited
 - More probabilistic approaches & weaker delivery possible




Publish-subscribe





Publish-subscribe

- Publish-subscribe
 - or distributed event systems
 - Working fundamentals
 - Publishers publish structured events to event service
 - Subscribers express interest in particular events
 - Event service
 - matches published events to subscriptions
 - Delivers suitable events
- 



Publish-subscribe


- Applications
 - Financial information systems
 - Live feeds of real-time data (including RSS)
 - Cooperative working
 - events of shared interest
 - Ubiquitous computing
 - location events, ... from infrastructure
 - Lots of monitoring applications


Stock trading system

- Let users see latest market prices of stock they care about
 - Info for a given stock arrives from multiple sources
 - Dealers only care about stocks they own (or might)
 - May only care to know above some threshold


Publish-subscribe

- Data provider
 - generate events (changes in stock value changes)
- Dealer process
 - creates subscription for each stock its user(s) express interest in





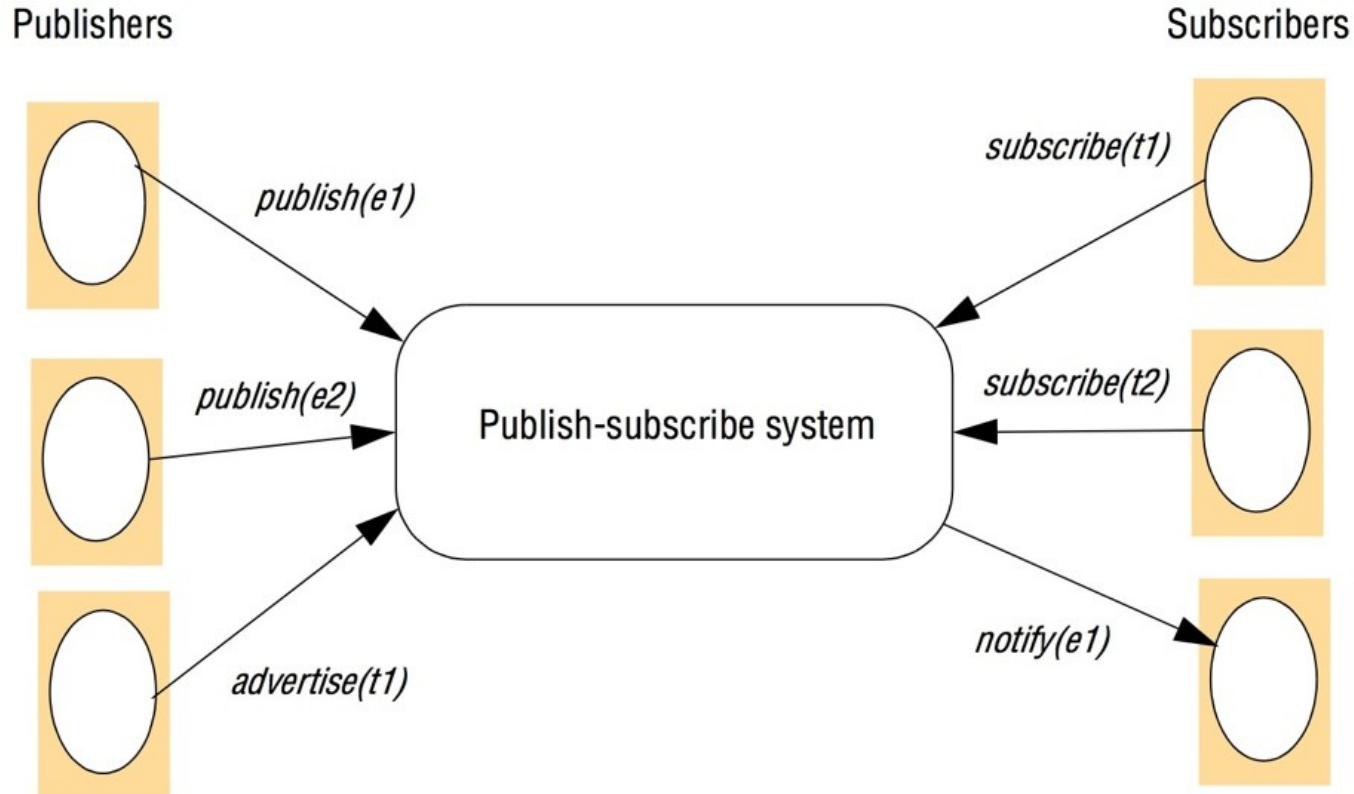
Publish-subscribe characteristics

- Heterogeneity
 - Able to glue together systems not designed to work together,
 - Have to come up with an external description of what can be subscribed to: simple flat, rich taxonomy, etc
 - Asynchronism
 - Decoupling means you never have to block
 - Possible delivery guarantees
 - All subscribers receive the same events (atomicity)
 - Events correctly delivered to subscribers at most once to subscribers (integrity)
 - message sent will be eventually delivered (validity)
 - Real-time
- 

Programming model

- Publishers
 - Disseminate event **e** through **publish(e)**
 - Register/advertise via a filter (pattern over all events):
 - **f: advertise (f)**
 - Expressiveness of pattern is the subscription model
 - Can also remove the offer to publish: **unadvertise (f)**
- Subscribers
 - Subscribe events that follow a filter/pattern
 - Receive events that match filter
 - Cancel their subscription:

Programming model



Subscription model


- Channel-based
 - Publishers publish to named channels
 - Subscribers get ALL events from channel
 - Very simplistic, no filtering (all other models below do)
- Topic-based or subject-based
 - Each notification expressed in multiple fields
 - one is the event topic
 - Subscribers choose topics
 - Allows hierarchical topics can help
 - e.g., old USENET rec.sports.cricket

Subscription model

- Content-based
 - Generalization of topic based
 - Subscription is expression over range of fields (constraints on values)
 - Far more expressive than channel-based or topic-based
- Type-based
 - Use object-based approaches with object types
 - Subscriptions defined in terms of types of events
 - Matching in terms of types or subtypes of filter
 - Ranges from coarse grained (type names) to fine grained (attributes and methods of object)
 - Advantage: clean integration with object-based programming languages



Publish-subscribe - Main concerns

- Deliver events efficiently to all subscribers that have filters that match the events
 - Security
 - Scalability
 - Failure handling
 - Quality of Service (QoS)
 - Tradeoffs:
 - Latency vs Reliability
 - Ease in implementation vs Expressiveness of events types
- 

Architecture

- Centralized schemes simple
 - Implementing channel-based or topic-based simple
 - Map channels/topics onto groups
 - Use the group's multicast (possibly reliable, ordered, ..)
 - Implementation of content/type/ more complicated
- Distributed
 - most implementations are network of brokers
- Some implementations are peer-to-peer (P2P)
 - All publisher and subscriber nodes act as the pub-sub broker

Publish-subscribe Distributed

