



Systems programming


13 - Synchronization

MEEC LEEC MEAer LEAer MEIC-A
João Nuno Silva





Bibliography

- Uresh Vahalia, Unix Internals the new frontier, chapter 7
 - Sun Micro systems, Multithreaded Programming Guide, chapter 4
 -
- 

Mutexes

- Used to guard Critical Regions
- Used to implement mutual exclusion
- Lock unlock
 - Should be done by the same task
- Simple interactions
 - Only one tasks running protected code



Other synchronization

- Multi-threaded programs may require different synchronization
 - Wait inside a critical region!!!!
 - Different resource guarding
 - Different synchronization patterns
 - Just thead execution coordination



Condition variables



Wait inside critical region

- Long critical regions
 - Serializes code
 - don't take advantage of concurrency
- Wait inside critical regions
 - For a certain condition/event
 - For a a specific variable value
- Impossible to do with regular mutexes

Wait inside critical region

- Thread 1

```
While (1){  
    lock(m1)  
    n - -  
    unlock(m1)  
}
```

Correct
But
active wait

- Thread 2

```
While (1){  
    lock(m1)  
    if (n == 0)  
        n = 100  
    unlock(m1)  
}
```

Condition Variables

- Programmers need additional mechanisms to
 - wait inside locked regions
- Without holding the lock while waiting
 - To allow other threads to enter the locked region
- Condition variables make it possible to
 - sleep inside a critical section
- Condition Variables
 - Atomically release the lock & go to sleep
 - Atomically acquire the lock and wake up




Condition Variables

- Each condition variable
 - Consists of a queue of threads
 - Provides three operations
 - Wait();
 - Atomically release the lock and go to sleep
 - Reacquire lock on return
 - Signal();
 - Wake up one waiting thread, if any
 - Broadcast();
 - Wake up all waiting threads
 - The queue contains threads that
 - Are sleeping/blocked
 - Will be waken by the signal/broadcast



Condition Variables

- Condition variable is associated with a mutex
 - That guards the critical region where thread wait
 - Wait
 - Releases and acquire corresponding mutex
 - Condition variables guarantee that
 - only one thread is inside the critical region
- 

An Example of Using Condition Variables

```
AddToQueue() {  
    lock(M1);  
    insert_item  
    unlock(M1);  
}  
RemoveFromQueue() {  
    while (1)  
        lock(M1);  
    if(something on queue)  
        break  
    else  
        unlock(m1)  
    remove_item  
    unlock(M1);  
    return item;  
}
```

```
AddToQueue() {  
    lock(M1);  
    insert_item  
    Signal(CV1, M1);  
    unlock(M1);  
}  
RemoveFromQueue() {  
    lock(M1);  
    Wait(CV1, M1);  
    remove_item  
    unlock(M1);  
    return item;  
}
```



Condition Variables

- Mutexes implement synchronization by controlling thread access to data
 - If before entering
- Condition variables allow threads to synchronize based upon the actual value of data.
 - If after entering

Condition Variables

- Without condition variables
 - The programmer would need to have threads continually polling to check if the condition is met.
 - Usually in a critical section
- A condition variable
 - Avoids polling (a.k.a. “busy wait”)
 - Useful when a thread needs to wait for a certain condition to be true.
 - Condition modified inside a critical region

Pthreads conditional var

- A condition variable is a `pthread_cond_t`

```
#include <pthread.h>
```

```
pthread_cond_t cond;
```

- Should be initialized before being used

```
int pthread_cond_init(pthread_cond_t *, const  
pthread_condattr_t *);
```

- 1º parâmetro: address of condition variable
 - 2º parâmetro: attributes (can be NULL)
- static initialization
 - `cond=PTHREAD_COND_INITIALIZER;`

Pthreads conditional var

- In pthreads, there are four relevant procedures involving condition variables:
 - `pthread_cond_init(pthread_cond_t *cv, NULL);`
 - `pthread_cond_destroy(pthread_cond_t *cv);`
 - `pthread_cond_wait(pthread_cond_t *cv, pthread_mutex_t *lock);`
 - `pthread_cond_signal(pthread_cond_t *cv);`



Conditional Vars - creation/destruction

- Condition variables must be declared with type **pthread_cond_t**, and must be initialized before they can be used.
 - Dynamically
 - `pthread_cond_init(cond, attr);`
- `pthread_cond_destroy(cond)`
 - used to free a condition variable that is no longer needed.

pthread_cond_wait

- `pthread_cond_wait(cv, lock)`
 - is called by a thread when it wants to block and wait for a condition to be true.
- It is assumed that the thread has locked the mutex indicated by the second parameter.
- The thread releases the mutex, and blocks until awakened by a `pthread_cond_signal()` call from another thread.
- When it is awakened,
 - it waits until it can acquire the mutex,
 - once acquired, it returns from the `pthread_cond_wait()` call.

pthread_cond_wait

- Waiting on a condition variable:

```
int pthread_cond_wait(  
    pthread_cond_t * cv,  
    pthread_mutex_t * mux);
```

- wait until other thread signals/broadcasts
- Guarantees that **mutex** is locked after return

• pthread_cond_signal

- pthread_cond_signal()
 - checks to see if there are any threads waiting on the specified condition variable.
 - If not, then it simply returns.
- If there are threads waiting,
 - then one is awakened.
- No assumption about the order in which threads awake
 - It is natural to assume that they will be awakened in the order in which they waited, but that may not be the case...
- Use loop or pthread_cond_broadcast() to awake all waiting threads.

Condition variables

- `pthread_timedwait`
- `pthread_cond_timedwait`
- `pthread_mutex_lock();`

• `while(condition_is_false)`

• `pthread_cond_wait();`

• `pthread_mutex_unlock();`

locked

locked

Signaling waiting thread

- the thread can change the condition variable
 - unlocks at least one thread
 - **`int pthread_cond_signal(pthread_cond_t *);`**
 - unlocks all threads
 - **`int pthread_cond_broadcast(pthread_cond_t *);`**
- Other Thread only resumes after this thread releases mutual exclusion

Condition variables

```
void *producer(void *) {  
    Produce_data()  
    pthread_mutex_lock(&data_mutex);  
    insert_data();  
    data_avail = 1;  
    pthread_cond_signal(&data_cond);  
  
    pthread_mutex_unlock(&data_mutex);  
}
```

```
void *consumer(void *) {  
    pthread_mutex_lock(&data_mutex);  
    while( !data_avail ) {  
        /* sleep on condition variable*/  
        pthread_cond_wait(&data_cond,  
                           &data_mutex);  
    }  
    /* woken up */  
    extract_data();  
    if (queue is empty)  
        data_avail = 0;  
    pthread_mutex_unlock(&data_mutex);  
}
```

Condition variables

task Producer

```
void *producer(void *) {  
    Produce_data()  
    pthread_mutex_lock(&data_mutex);  
  
    insert_data();  
    data_avail = 1;  
    pthread_cond_signal(&data_cond);  
    pthread_mutex_unlock(&data_mutex);  
}
```

task Consumer 1

```
void *consumer(void *) {  
    pthread_mutex_lock(&data_mutex);  
    while( !data_avail ) {  
        /* sleep on cond var*/  
        pthread_cond_wait(&data_cond, &data_mutex);  
    }  
    extract_data();  
    if (queue is empty)  
        data_avail = 0;  
    pthread_mutex_unlock(&data_mutex);  
}
```

Condition variables

- Loop around wait
 - Other threads may be woken up first.
 - It is possible that another thread might acquire the mutex first and change the state
 - Designing for “loose” predicates may be simpler.
 - Define on condition variables that indicate possibility rather than certainty.
 - signaling a condition variable would mean
 - there may be something” for the signaled thread to do, rather than “there is something” to do.
 - Spurious wake-ups can occur.
 - On some implementations, a thread waiting on a condition variable may be woken up even though no other thread actually signaled the condition variable.
 - Due to optimizations and explicitly permitted by SUSv3.

Read-Write Locks

Read-Write Locks

- Critical regions can be relaxed
- Some of the operations on the resource
 - Can be executed concurrently (just read)
 - Others should be serialized (writes)
- Allow concurrent reads speeds program

Read-Write Lock

- Read-write locks permit concurrent reads and exclusive writes to a protected shared resource.
- The read-write lock is a single entity that can be locked in read or write mode.
- To modify a resource, a thread must first acquire the exclusive write lock.
 - An exclusive write lock is not permitted until all read locks have been released.
-
- Read-write locks support concurrent reads of data structures
 - read operation does not change the record's information.
- When the data is to be updated
 - the write operation must acquire an exclusive write lock.

Not fair / Fair scheduling

- Th1 ReadLock (enter)
- Th2 ReadLock (enter)
- Th3 WriteLock (blocked)
- Th2 unlock (leave)
- Th4 ReadLock (enter)
- Th1 unlock (leave)
- Th2 ReadLock (enter)
-
- Th3 is in starvation

- Th1 ReadLock (enter)
- Th2 ReadLock (enter)
- Th3 WriteLock (blocked)
- Th2 unlock (leave)
- Th4 ReadLock (blocked)
- Th1 unlock (leave)
- Th3 enters
- Th2 ReadLock (blocked)
- .

Read-Write Lock

- Initialization

- `int pthread_rwlock_init(pthread_rwlock_t *restrict rwlock,`
- `const pthread_rwlockattr_t *restrict attr);`
 - `rwlock` will contain the reference to the lock
 - `attr` can be `NULL`

- Destruction

- `int pthread_rwlock_destroy(pthread_rwlock_t *rwlock);`

- locking

- `int pthread_rwlock_rdlock(pthread_rwlock_t *rwlock);`
- `int pthread_rwlock_wrlock(pthread_rwlock_t *rwlock);`

- unlocking

- `int pthread_rwlock_unlock (pthread_rwlock_t *rwlock);`

Sempahores

Mutex

- Mutex is rigid
 - Count only has 2 values – 0 1
 - Order of operations is fixed: lock → unlock
 - Only owner can unlock
- Some problems require relaxation of these properties
 - “Critical regions” with multiple threads
 - One thread unblocking another

Semaphores

- Abstraction for a counter and a task list, used for synchronization purposes.
 - Proposed by Dijkstra in 1965.
 - Do not require active wait.
 - All modern OS include a version of semaphores (Unix, Windows, ...)
- `typedef struct{`
 - `unsigned counter;`
 - `processList *queue;`
 - `}sem_t;`

Semaphores

- S.counter
 - Defines how many tasks can pass the semaphore without blocking
- If s.counter = 0
 - The next task to try to enter will get blocked
- If s.counter is always ≤ 1
 - Mutual exclusion is guaranteed
 - The length of the queue depends on the number of tasks waiting to enter

Semaphores

- `down(S)`
 - Used by a tasks when trying to access a resource
 - Access is given by another task issuing `up(S)`.
 - Task can get blocked if capacity is full
 - `counter == 0`
- `up(S)`
 - Used by a task to signal the resource availability
 - `up(S)` is not blocking

Semaphores

- Wait(S), down(S), or P(S)
 - if (S.counter>0)
 - S.counter--;
 - else
 - Block(S); /* insert S into the queue */
- Signal(S), up(S), or V(S)
 - if (S.queue!=NULL)
 - WakeUp(S); /* removes a process from queue */
 - else
 - S.counter++;
- Wakeup and block depend on the OS
- P(S) e V(S) come from dutch words prolaag (decrement) and verhoog (incrementar)

Semaphores

- Semaphores allow implementation of other mechanisms
-
- Critical region
 - Initialize a semaphore with counter = 1
 - Do a down when entering the critical region
 - Do a Up when leaving the critical region
- Rendezvous
 - Initialize a semaphore with counter = 0
 - The tasks that does down(S) will get blocked
 - Until other task does the UP(S)
 - Two tasks reandez-vous and continued together

POSIX Semaphores

- A semaphore is an integer whose value is never allowed to fall below zero.
- POSIX includes a set of functions
 - `man sem_overview`
 - `#include <semaphore.h>`
 - Link program with `-lpthread`

POSIX Semaphores

- Two operations can be performed on semaphores:
 - increment the semaphore value by one
 - `sem_post()`
 - and decrement the semaphore value by one
 - `sem_wait()`
- If the value of a semaphore is currently zero,
 - then a `sem_wait` operation will block until the value becomes greater than zero.

POSIX Semaphores

- POSIX offer two forms of semaphores with respect to creation
 - Named semaphores
 - Identified by a global name (null terminated string started with /
 - Multiple processes can operate on the same named semaphore by passing the same name to `sem_open`
 - Unnamed (memory-based semaphores)
 - Created in memory shared by multiple threads

POSIX Semaphores

- POSIX offers two sharing mechanism
 - Not shared
 - Just threads of the process
 - Shared
 - Unnamed - Shared by threads in related processes parent children or using shared memory
 - Named - Shared by threads in multiple processes
 - Semaphores are also classified by the maximum number of tasks (N) that can access the resources
 - if N equals 1 => binary semaphore

POSIX Semaphores

Named		unnamed
sem_open()		sem_init()
	sem_wait() sem_trywait() sem_post() sem_getvalue()	
sem_close() sem_unlink()		sem_destroy()

Semaphore

up

down

POSIX

sem_post

sem_wait

POSIX unnamed Semaphores

- An unnamed semaphore is a variable of type **sem_t**
 - `#include <semaphore.h>`
 - `sem_t sem;`
- Should be initialized before being used by
 - Child processes,
 - threads.
- **int sem_init(sem_t *sem, int pshared, unsigned int value)**
 - `sem_init()` initializes the unnamed semaphore at the address pointed to by **sem**.
 - The **pshared** argument indicates whether this semaphore is to be shared between the threads of a process, or between processes.
 - The **value** argument specifies the initial value for the semaphore.

POSIX unnamed Semaphores

- Unused semaphores should be destroyed
 - `int sem_destroy(sem_t *)`;

```
sem_t semaforo;  
if (sem_init(&semaforo,0,1)==-1)  
    perror("Falha na inicializacao");  
  
if (sem_destroy(&semaforo)==-1)  
    perror("Falha na eliminacao");
```



POSIX named Semaphores

- Allows the synchronization of processes without shared memory
 - Use a global name
 - string following the form **/name**
- Named semaphores are installed/located in `/dev/shm`, with the name **sem.name**
- A semaphore should be opened before used

POSIX named Semaphores

- `sem_t *sem_open(const char *name, int oflag)`
- `sem_t *sem_open(const char *name, int oflag, mode_t mode, unsigned int value)`
 - `sem_open()` creates a new POSIX semaphore or opens an existing semaphore.
 - The semaphore is identified by name.
 - The `oflag` argument specifies flags that control the operation of the call.
 - If `O_CREAT` is specified in `oflag`, then the semaphore is created if it does not already exist.
- If `O_CREAT` is specified in **`oflag`**, two additional arguments must be supplied
 - **`mode_t`**, permissions (owner, group, user)
 - The value argument specifies the initial value for the new semaphore..
- If **`oflag`** contains `O_CREAT` and `O_EXCL`
 - Error if semaphore already exists
- If **`oflag`** is `O_CREAT` and semaphore exists
 - 3rd and 4th parameters are ignored

POSIX named Semaphores

- When the semaphore is of no use should be closed
 - `int sem_close(sem_t *);`
- The last process should
 - Close the semaphore (`sem_close`)
 - Remove the corresponding file
 - `int sem_unlink(const char *);`
- If a process maintains the semaphore opened
 - `sem_unlink` is blocked until the semaphore is closed
- If the semaphore is not closed/unlinked
 - New uses of the semaphore are undefined....

POSIX Semaphores


- down(S) is implemented by the function
 - `int sem_wait(sem_t *);`
 - If counter is zero
 - The thread executing the function is blocked.
 - Remaining thread in the process continue executing
- up(S) is implemented by
 - `int sem_post(sem_t *);`



Barriers / rendezvous

Barriers

- Synchronization mechanism that blocks threads until a defined number of threads arrives at the barrier



- Are used when processing is done in steps, that require the completion of a number of threads
 - The next step is only started (by all threads at the same time) after all threads have finished the preceding step

Barriers

- Pthreads provides the type `pthread_barrier_t`
**`pthread_barrier_init(pthread_barrier_t *,`
**`pthread_barrierattr_t *,`
`unsigned int)`****

- 1st parameter - barrier object
- 2nd parameter - barrier attributes
- 3rd parameter - number of threads that must call `pthread_barrier_wait()` before any of them successfully return from the call.

`pthread_barrier_destroy(pthread_barrier_t *)`

Barriers

- Wait / Synchronization

```
int pthread_barrier_wait(pthread_barrier_t *)
```

- All threads block
- When the count value (3rd argument from ini) is reached
 - Count thread unblock and start executing
 - Function returns
- Return value
 - one thread - PTHREAD_BARRIER_SERIAL_THREAD,
 - All others - 0