



# Systems programming


## 12 - Shared data

MEEC LEEC MEAer LEAer MEIC-A  
João Nuno Silva





# Bibliography

- Sun Microsystems, Multithreaded Programming Guide, chapter 4
  - 
  - Michael Kerrisk, The Linux Programming Interface, Chapter 30
- 



# Shared data in threads

- Shared variables
  - Global variable in C
  - Memory region allocated
- Shared data should be avoided
- But sometimes is required
  - Two thread access (modify) same variable

# Data access

- Access to data in memory is not atomic
  - It is necessary to load, process and store
  - Even without caches
  - Even on a single computer
- The simple C operation (`i++`) is complex

General	MIPS	x86	ARM
Load Reg $\leftarrow$ Mem	lw \$t0, label	Mov ax, label	LDR R2, [label]
Increment register	addiu \$t0, \$t0, 1	Inc ax	ADDS R2, R2, 1
Store Mem $\leftarrow$ reg	sw \$t0, label	Mov, label, ax	STR R2, [label]

# Instruction ordering

- Single-threaded application
  - Order of instructions is the order of C Code
- Multi-threaded application
  - Order of instructions is one of the possible combinations

Main	Thread 1	Thread 2
N = 0		
pthread_create		
	Load R1 $\leftarrow$ N R1 $\leftarrow$ R1 + 1 Store N $\leftarrow$ R1	Load R1 $\leftarrow$ N R1 $\leftarrow$ R1 + 1 Store N $\leftarrow$ R1




# Thread synchronization

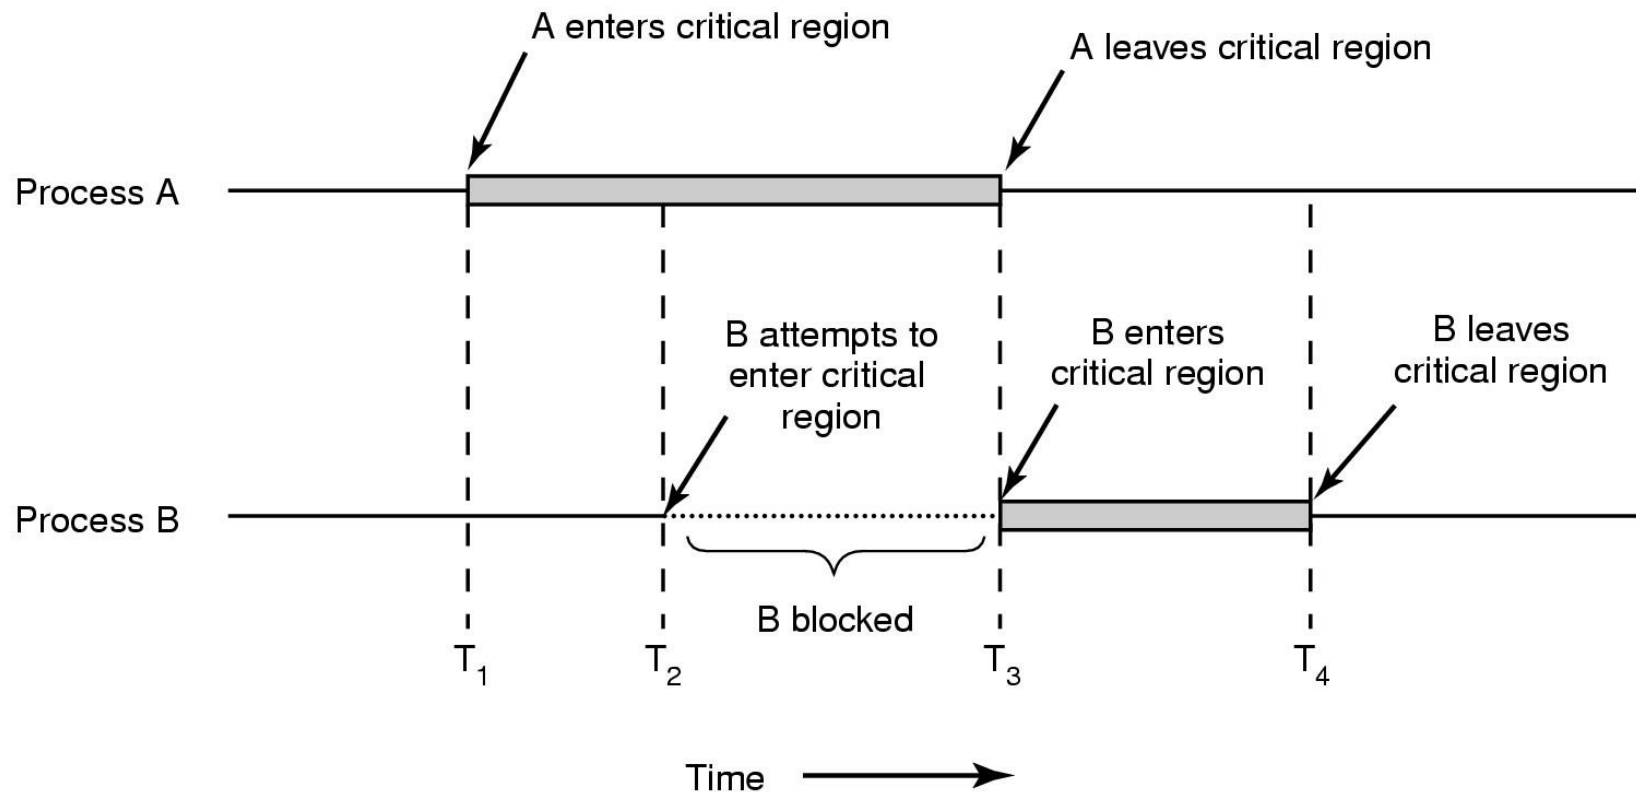
- Threads execution order
  - defined by the programmer not efficient
  - Multiple combinations
- Protect accesses to shared variables
  - Define critical region
  - Only one such critical region is executed at a time
    - Only one thread is inside such section
- Implement mutual exclusion

# Critical region

- Piece of code where resources are shared
  - That can should be executed by one tasks at a time
  - Is delimited by read/write instructions to the shared resource
- If a task is inside a critical region
  - Other task trying to enter should be blocked



# Critical region





# Critical region

- Requirements to be satisfied
  - Mutual exclusion
    - Only one task can be inside the Critical region
  - Progress
    - A task inside the Critical region can not block other task from entering
  - Limited wait
    - A task should wait a limited amount of time before entering
- A task should remain inside the Critical Region for a limited time
- The speed and number of processor is undefined
  - Should work for any combination of CPU/tasks

# Mutual exclusion

- Mechanism that assures that at most one task is inside a critical region
  - No more than one task is executing the critical region code
- All other tasks are
  - Running non critical code
  - Blocked
- Consequences of mutual exclusion
  - starvation (live-lock)
    - A task is able to be executed, but is never scheduled
  - deadlock
    - Due to coding problem, several tasks are waiting for being unblocked by other tasks that are in the same state

# Critical region / Mutual exclusion

```
do {
```

```
    non_critical_code()
```

```
    EnterRegion()
```

**zero or more  
tasks**

```
    CriticalRegion(shared_data)
```

**Zero or one  
tasks**

```
    LeaveRegion()
```

```
    non_critical_code()
```

```
} while(X);
```

**All other tasks**

# Locks

- The simplest mechanism to ensure
  - mutual exclusion of critical sections

- Spin locks
- Busy waiting is inefficient
  - NOPs use energy
- Requires special instructions
  - While is a critical region
  - Test-and-set
  -

```
while (lock == 1) {}  
lock = 1;  
// enter crit reg  
criticalRegion()  
lock = 0;  
// leave crit reg
```



# Mutual exclusive locks - mutex

- OS object that
  - have 2 values (0 or 1)
  - Have a tasks waiting list
  - Is managed by the operation system
    - To suspend and run tasks
- A Mutex can be in the following states
  - Unlocked - no tasks inside the critical region
  - Locked by a single task - the one in the critical region

# Mutex usage

- After defining a critical region
  - Assign it a mutex
  - Initialize a mutex
- Each thread
  - Call `mutex_lock` when entering the critical region
  - Call `mutex_unlock` when leaving the critical region
- `Mutex_lock`
  - Tasks can be placed in a list, waiting for a `mutex_unlock`
- `Mutex_unlock`
  - Can only be called by the thread inside the critical region

# Mutex usage

- Good programming
  - `mutex_lock(m1)`
  - `CriticalRegion()`
  - `mutex_unlock(m1)`
  - Small critical region
  - Same code on all tasks
- Bad programming
  - Long critical regions
  - `mutex_lock(m1)`
    - Inside critical region
    - Without unlock
  - `mutex_unlock()`
    - Not owner
    - Outside Critical Region

# POSIX mutexes

- POSIX mutexes are associated to Pthreads:
  - include file **#include <pthread.h>**
  - Compile with **-lpthread**
- The corresponding data type is
  - **pthread\_mutex\_t mux;**
- A mutex should be initialized before used
  - **mutex=PTHREAD\_MUTEX\_INITIALIZER;**
- A mutex is destroyed by calling:
  - **int pthread\_mutex\_destroy(pthread\_mutex\_t \*);**



# POSIX mutexes

- mutex locking
  - `int pthread_mutex_lock(pthread_mutex_t *mutex);`
  - Blocks thread until resource is available/can enter critical region
  - Returns when task enters critical region
  - Returns 0 in case of success
- Mutexes should be locked for the minimum amount of time
- Mutex unlock
  - `int pthread_mutex_unlock(pthread_mutex_t *mutex);`
  - Returns 0 in case of success
  - Allow other thread to enter region
  - Unblock a thread from `pthread_mutex_lock`

# POSIX Spin Locks

- Spin locks are a low-level synchronization mechanism
  - suitable for tiny critical regions
- When the calling thread requests a spin lock that is already held by another thread
  - the second thread spins in a loop to test if the lock has become available.
- When the lock is obtained it should be held only for a short time,
  - as the spinning wastes processor cycles.

# POSIX Spin Locks

- POSIX spin locks are associated to Pthreads:
  - include file `#include <pthread.h>`
  - Compile with `-lpthread`
- The corresponding data type is
  - `pthread_spinlock_t *lock;`
- initialization
  - `int pthread_spin_init(pthread_spinlock_t *lock, int pshared);`
- Destruction
  - `int pthread_spin_destroy(pthread_spinlock_t *lock);`



# POSIX Spin Locks

- locking
  - `int pthread_spin_lock(pthread_spinlock_t *lock);`
- unlocking
  - `int pthread_spin_unlock(pthread_spinlock_t *lock);`

# POSIX Spin Locks

- A programmer must be exceptionally careful with
  - the code, system configuration, thread placement, and priority
- If a thread creates deadlock while holding a spin lock
  - It will spin forever consuming CPU time.
- If a thread is scheduled off the CPU while it holds a spin lock
  - other threads will waste time spinning on the lock until
    - the lock holder is once more rescheduled
    - and releases the lock