



TÉCNICO
LISBOA

OPTIMIZATION & ALGORITHMS

MEEC

Project Report

Group: 4

Alexandre Leal (103098)

alexandre.b.leal@tecnico.ulisboa.pt

Diogo Ramos (100299)

diogo.ramos@tecnico.ulisboa.pt

Diogo Sampaio (103068)

diogo.sampaio@tecnico.ulisboa.pt

Francisco Tavares (103402)

francisco.carreira.tavares@tecnico.ulisboa.pt

2024/2025 – 1º Semestre, P1

1 Task 1: Theoretical Task

In this task, we demonstrate that the function

$$f_D(w_0, w) = \frac{1}{N} \sum_{n=1}^N 1_{R_-}(y_n C_{w_0, w}(x_n))$$

is not convex in the simplified case where $N = 1$ and $D = 1$. For $N = 1$, it simplifies to:

$$f_D(w_0, w) = 1_{R_-}(y_n C_{w_0, w}(x_n)).$$

We define two functions:

$$g(u) = 1_{R_-}(u), \quad f_2(w_0, w) = y_n C_{w_0, w}(x_n).$$

Thus, we express f_D as the composition:

$$f_D(w_0, w) = g \circ f_2.$$

If $g(u)$ is not convex, then f_D is also not convex if f_2 maps to a region where $g(u)$ is not convex. Observing figure 1, we can see that $g(u)$ is not convex.

To demonstrate, using the definition of convexity consider $x = -1$, $y = 1$, and $\alpha = 0.25$:

$$\begin{aligned} g((1 - \alpha)x + \alpha y) &\leq (1 - \alpha)g(x) + \alpha g(y) \\ \iff g(0.75 \cdot (-1) + 0.25 \cdot 1) &\leq 0.75g(-1) + 0.25g(1) \\ \iff g(-0.5) &\leq 0.75 \cdot 1 + 0.25 \cdot 0 \\ \iff 1 &\leq 0.75. \end{aligned}$$

This is false; hence, $g(u)$ is not convex, and so f_D is also not convex.

Note: The function f_2 maps \mathbb{R} to $\{+1, -1\}$ since y_n and $C_{w_0, w}$ belong to $\{+1, -1\}$, validating the points used in the example.




Figure 1: Plot of the function $g(u)$

2 Task 2. [Theoretical Task]

In this task we were asked to show that the function 1_{R_-} is majorized by h , and that $1_{R_-} \leq h(u)$ holds for all $u \in \mathbb{R}$. Furthermore we showed that h is a convex function.

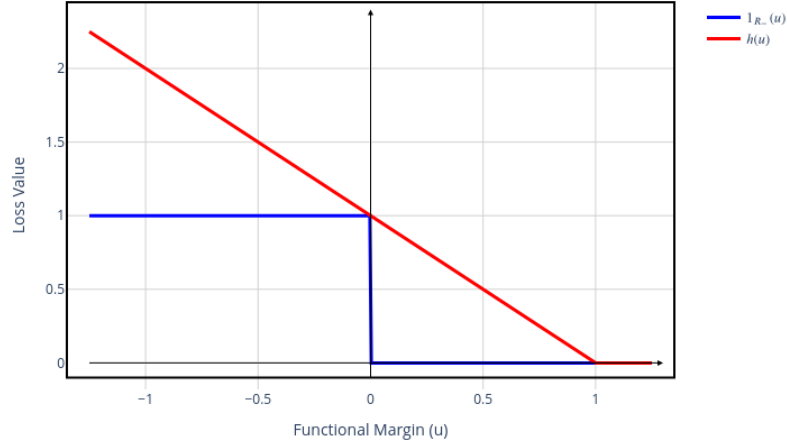


Figure 2: Plot of the indicator function $1_{R_-}(u)$ & hinge loss function $h(u)$

We want to prove that $1_{R_-}(u) \leq h(u)$ for all $u \in \mathbb{R}$. We can describe $1_{R_-}(u)$, and $h(u)$ as:

$$h(u) = \begin{cases} 0, & u \leq 1 \\ 1 - u, & u > 1 \end{cases}$$

$$1_{R_-}(u) = \begin{cases} 1, & u < 0 \\ 0, & u \geq 0 \end{cases}$$

For $u < 0$:

$$\begin{aligned} 1_{R_-}(u) \leq h(u) &\iff 1 \leq 1 - u \\ &\iff 0 \leq -u \iff u \leq 0 \quad \text{True} \end{aligned}$$

For $u \in [0; 1[$:

$$\begin{aligned} 1_{R_-}(u) \leq h(u) &\iff 0 \leq 1 - u \\ &\iff u \leq 1 \quad \text{True} \end{aligned}$$

For $u \geq 1$:

$$1_{R_-}(u) \leq h(u) \iff 0 \leq 0 \quad \text{True}$$

To show that h is a convex function we can observe figure 3, it's easy to see that no matter which 2 points in h we choose, if we draw a straight line between them, h will be always lower or equal then any point in that line, therefore, h is convex.

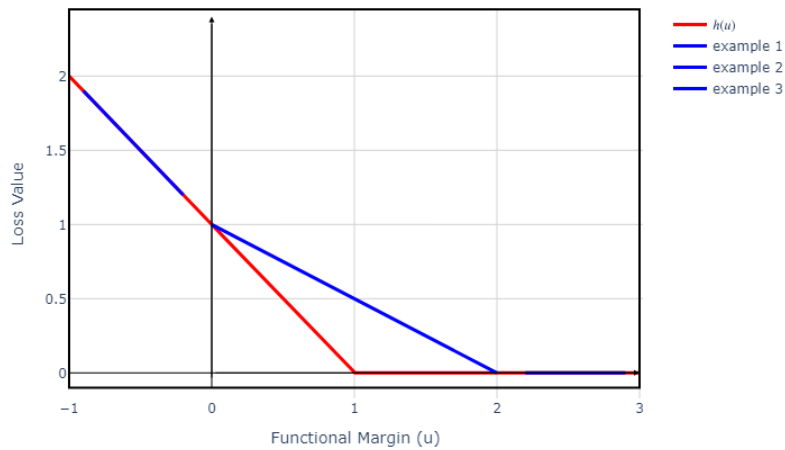


Figure 3: Plot of the hinge loss function $h(u)$ and some examples to show convexity

3 Task 3: Theoretical Task

We are asked to show that the function

$$g_D(\omega_0, \omega) = \frac{1}{N} \sum_{n=1}^N h(y_n(\omega_0 + x_n^T \omega))$$

is convex for any N and D .

The function $g_D(\omega_0, \omega)$ can be written as a sum of convex functions in the form:

$$g_D(\omega_0, \omega) = \frac{1}{N} \sum_{n=1}^N h(y_n(\omega_0 + x_n^T \omega)) = \frac{1}{N} h(y_1(\omega_0 + x_1^T \omega)) + \cdots + \frac{1}{N} h(y_N(\omega_0 + x_N^T \omega)).$$

If f_1, \dots, f_N are convex, then $g_D(\omega_0, \omega)$ is convex because a sum of convex functions is convex [slide 12 of module 3 of the theoretical slides]. Since all f_n have the same form, we only need to prove that f_n is convex for all $n \in \{1, \dots, N\}$.

We define

$$f_n = h(y_n(\omega_0 + x_n^T \omega)),$$

which can be written as $f_n = h(g(\omega_0, \omega))$. To show f_n is convex, we need to check two conditions:

1. $h(u)$ is convex, and 2. $g(\omega_0, \omega)$ is affine [slide 12 of module 3 of the theoretical slides].

First, $g(\omega_0, \omega)$ can be written as

$$g(\omega_0, \omega) = y_n(\omega_0 + x_n^T \omega) = \begin{bmatrix} y_n & y_n x_n^T \end{bmatrix} \begin{bmatrix} \omega_0 \\ \omega \end{bmatrix},$$

where

$$A = \begin{bmatrix} y_n & y_n x_n^T \end{bmatrix} \in \mathbb{R}^{1 \times (D+1)}, \quad b = 0, \quad \text{and} \quad x = \begin{bmatrix} \omega_0 \\ \omega \end{bmatrix} \in \mathbb{R}^{D+1}.$$

Thus, $g(\omega_0, \omega)$ is an affine map of the form $g(\omega_0, \omega) = Ax + b$, where A is a matrix and b is a vector.

Since affine functions preserve convexity, and we are given that $h(u)$ is convex (as shown in Task 2), it follows that $f_n = h(g(\omega_0, \omega))$ is convex. Therefore, $g_D(\omega_0, \omega)$ is convex.

4 Task 4. [Theoretical Task]

We can write g as $g_d(w_0, w) + p \cdot (\|w\|_2)^2$, where $p > 0$, which is in the form $t_1 f_1 + t_2 f_2$. Therefore, if g_d is convex (proven in Task 3) and $(\|w\|_2)^2$ is convex, then g is convex.

We can write $(\|w\|_2)^2$ as $f \circ h(w)$, with $h(w) = \|w\|_2$ and $f(u) = (u_+)^2$. We can say this because for $u \geq 0$, $u_+ = u$ and we know that $h(w) \geq 0$ for all w . $f \circ h(w)$ is convex if f is convex and non-decreasing and h is convex. h is an elementary convex function (the l_p norm) [slide 7 of module 3 of the theoretical slides], so it is convex.

$f(u)$ can be described as:

$$\begin{cases} f(u) = u^2, u \geq 0 \\ f(u) = 0, u < 0 \end{cases} \quad (1)$$

Because f is continuous, if both segments are non-decreasing, then f is also non-decreasing. For $u < 0$, f is constant, so it is non-decreasing. For $u \geq 0$, $f'(u) = 2u$. Since $u \geq 0$, it means the slope of f is always positive, which implies that f is non-decreasing. This means $(\|w\|_2)^2$ is convex and therefore g is also convex.

5 Task 5. [Theoretical Task]

To determine if the function g is strongly convex we will evaluate the case where $w = 0$ and $N = 1$. In this range of values we get $g(\omega_0, \omega) = h(y\omega_0)$. If we analyze the range where $y\omega_0 \geq 1$, g is a null constant, which means it is not strictly nor strongly convex. Equivalently, if we consider $f = g - \frac{-m\|\omega_0\|^2}{2} = \frac{-m\|\omega_0\|^2}{2}$, in the same range, for all positive values of m , f is clearly not convex, which equates to g not being strongly convex.

Now using the definition of strongly convex, for g to be SCVX,

$$g((1 - \alpha)x + \alpha y) \leq (1 - \alpha)g(x) + \alpha g(y) - \frac{n\alpha(1 - \alpha)}{2} \|x - y\|^2 \quad (2)$$

Let's take for example $x = 20$ and $y = 30$ and $\alpha = 0.5$,

$$g(10 + 15) \leq 0.5g(20) + 0.5g(30) - \frac{n}{8} \|x - y\|^2 \quad (3)$$

Knowing that for $x > 1$, $g(x) = 0$, $x = 20, y = 30$

$$0 \leq -n \frac{\|x - y\|^2}{8} \quad \frac{\|x - y\|^2}{8} > 0 \quad (4)$$

$0 \leq -n$ is false for all positive values of n therefore g is not strongly convex.

Furthermore, looking at the graph of the hinge loss function at the end of Task 2, we can also see the constant part of the function mentioned above, proving that g is not strongly convex.

6 Task 6. [Numerical Task]

In this task, we solve the following optimization problem numerically, using CVX in Python. The problem involves classifying images of handwritten digits from the MNIST dataset, where the digits 0 and 1 are considered. We use a linear classifier based on the hinge loss function with regularization, as stated in equation (5) of the problem.

The optimization problem we are solving is defined as follows:

$$\min_{w_0, w} \left[\frac{1}{N} \sum_{n=1}^N \max(0, 1 - y_n(w^T x_n + w_0)) + \lambda \|w\|_2^2 \right] \quad (5)$$

The first term in the objective function represents the hinge loss, and the second term is an ℓ_2 regularization term to prevent overfitting.

After solving the optimization problem, the following optimal parameters were obtained:

$$w_0 \approx 0.255$$

$$w \approx [0, 0, 0, \dots, -1.76 \times 10^{-10}, -2.65 \times 10^{-10}, \dots, 2.43 \times 10^{-3}, \dots]$$

We also evaluated the classifier error rate f_D on both the training and test datasets using the function:

$$f_D = \frac{1}{N} \sum_{n=1}^N \mathbf{1}(y_n \neq \text{sign}(w^T x_n + w_0))$$

The error rates were as follows for λ :

- Training dataset error rate: 0.00%
- Test dataset error rate: 0.12%

In addition to that, we also changed λ to 0.5 to verify our code and noticed, as expected, that f_d evaluates to 0.25% for the training set and 0.25% for the test set.

7 Task 7. [Theoretical Task]

7.1 Task 7 a.

This task is divided into two parts. First, we are asked to show that

$$\tilde{x} = x - P \operatorname{sgn}(yw)$$

solves the following equation:

$$\begin{aligned} & \underset{\tilde{x}}{\text{minimize}} \quad y(w_0 + \tilde{x}^T w) \\ & \text{subject to} \quad |\tilde{x}_d - x_d| \leq P, \quad \text{for } 1 \leq d \leq D. \end{aligned} \tag{6}$$

To solve this task, we started by expanding the cost function:

$$f(\tilde{x}) = y(\omega_0 + \tilde{x}^T \omega) = y\omega_0 + y\tilde{x}^T \omega,$$

Since $y\omega_0$ is a constant, we wish to minimize $y \sum_{d=1}^D \tilde{x}_d \omega_d$ where $\tilde{x}_d \in [x_d - P; x_d + P]$.

- If $y\omega_d > 0$, we should pick the smallest \tilde{x}_d , $\tilde{x}_d = x_d - P$.
- If $y\omega_d < 0$, we should pick the largest \tilde{x}_d , $\tilde{x}_d = x_d + P$.
- If $y\omega_d = 0$, any value of ω_d works because $y\tilde{x}_d \omega_d = 0$

$$\tilde{x}_d = \begin{cases} x_d - P, & \text{if } y\omega_d > 0 \\ x_d + P, & \text{if } y\omega_d < 0 \\ \text{any value,} & \text{if } y\omega_d = 0 \end{cases}$$

From this analysis, we conclude that the optimal solution is given by $\tilde{x} = x - P \operatorname{sgn}(yw)$. Next, we need to verify whether this solution satisfies the specified constraints:

$$\begin{aligned} |\tilde{x}_d - x_d| &\leq P, \quad d \in \{1, 2, 3\} \\ |x - P \operatorname{sgn}(yw) - x| &= |-P \operatorname{sgn}(yw)| = P \leq P \end{aligned} \tag{7}$$

This confirms that $\tilde{x} = x - P \operatorname{sgn}(yw)$ solves the optimization problem stated in (1).

7.2 Task 7 b.

In the second part of the task, we are asked to show that for $\tilde{x} = x - P \operatorname{sgn}(yw)$, the cost function in (1) evaluates to:

$$y(\omega_0 + x^T \omega) - P \|yw\|_1.$$

Substituting \tilde{x} into the cost function:

$$\begin{aligned} f(\tilde{x}) &= y (w_0 + (x - P \operatorname{sgn}(yw))^T w) \\ &\iff y (w_0 + x^T w - P \operatorname{sgn}(yw)^T w) \\ &\iff y \left(w_0 + x^T w - P \sum_{d=1}^D \operatorname{sgn}(yw_d) w_d \right) \\ &\iff y(w_0 + x^T w) - P \sum_{d=1}^D |yw_d| \\ &\iff y(w_0 + x^T w) - P \|yw\|_1. \end{aligned}$$

which confirms the evaluation of the cost function.

Thus, we conclude that the derived expression correctly evaluates the cost function for $\tilde{x} = x - P \operatorname{sgn}(yw)$.

8 Task 8. [Numerical Task]

Using the result of Task 7, we replaced (x_n, y_n) by (\tilde{x}_n, y_n) with $\tilde{x}_n = x - P \operatorname{sgn}(yw)$ and evaluated the function f_d defined in (2), obtaining the classifier of Task 6's error rate on the attacked test dataset. With $\lambda = 0.1$, as used in task 6, and $P=0.18$, f_d evaluates 43.56% . We also changed λ to 0.5 to verify our code and noticed, as expected, that f_d evaluates 21.88%, which is very close to the expected value 21.9%

9 Task 9. [Numerical Task]

In this section we want to solve the problem of minimizing the effect of an attack on our classifier. This corresponds to solving the problem defined as:

$$\min_{w_0, w} \left[\frac{1}{N} \sum_{n=1}^N \max(0, 1 - (y_n(w^T x_n + w_0) - P||y_n w||_1)) + \lambda ||w||_2^2 \right] \quad (8)$$

with $P = 0.18$ and $\lambda = 0.1$.

On the training dataset f_D evaluates to 0.75% and in the test dataset f_D evaluates to 0.44%. As for the attacked dataset, f_D evaluates to 2.19%.

We can observe that, while in the test and training sets, the classifier performs slightly worse, which is expected since it is not the ideal classifier for unaltered data, we see a large improvement on the performance on the attacked dataset, since this is the loss function we strive to minimize with this formulation.

10 Task 10. [Numerical Task]

In Task 10, we aim to solve the problem of fitting a piecewise-linear signal to a set of noisy measurements. The signal is modeled as a weighted combination of linear models, and we need to optimize the parameters of these models to minimize the fitting error.

The goal is to minimize the following objective function f :

$$\min_{s_1, r_1, \dots, s_K, r_K, u_1, v_1, \dots, u_{K-1}, v_{K-1}} \sum_{n=1}^N (\hat{y}(x_n) - y_n)^2 \quad (9)$$

where $\hat{y}(x_n)$ is the predicted value at time x_n , given by:

$$\hat{y}(x_n) = \sum_{k=1}^K w_k(x_n) \hat{y}_k(x_n) \quad (10)$$

and y_n represents the measured values.

Each linear model $\hat{y}_k(x)$ is defined as:

$$\hat{y}_k(x) = s_k x + r_k \quad (11)$$

where s_k is the slope and r_k is the intercept for the k -th model.

The weights $w_k(x)$, which determine how the models are combined, are given by the softmax function:

$$w_k(x) = \frac{e^{u_k x + v_k}}{\sum_{j=1}^K e^{u_j x + v_j}} \quad (12)$$

These weights $w_k(x)$ ensure that the combination of linear models is smooth across the different regions of the signal.

When solving for the gradient of the objective function, we first get:

$$\nabla f = 2 \sum_{n=1}^N (\hat{y}(x_n) - y_n) \cdot \nabla \hat{y}(x_n) \quad (13)$$

where $\nabla \hat{y}(x_n)$ represents the Jacobian matrix J .

The Jacobian matrix J consists of the partial derivatives of $\hat{y}(x_n)$ with respect to the variables s_k, r_k, u_k , and v_k :

$$J_{n,p} = \frac{\partial \hat{y}(x_n)}{\partial p} \quad (14)$$

where p can represent s_k, r_k, u_k , or v_k .

To compute the partial derivatives, we need to evaluate the derivatives of $w_k(x)$ and $w_k(j)$:

1. Derivative of $w_k(x)$ with respect to u_k :

$$\begin{aligned}
\frac{\partial w_k(x)}{\partial u_k} &= \frac{\partial}{\partial u_k} \left(\frac{e^{u_k x + v_k}}{\sum_{i=1}^K e^{u_i x + v_i}} \right) \\
&= \frac{\frac{\partial}{\partial u_k} (e^{u_k x + v_k}) \cdot \left(\sum_{i=1}^K e^{u_i x + v_i} \right) - e^{u_k x + v_k} \cdot \frac{\partial}{\partial u_k} \left(\sum_{i=1}^K e^{u_i x + v_i} \right)}{\left(\sum_{i=1}^K e^{u_i x + v_i} \right)^2} \\
&= \frac{e^{u_k x + v_k} \cdot x \cdot \left(\sum_{i=1}^K e^{u_i x + v_i} \right) - e^{u_k x + v_k} \cdot (e^{u_k x + v_k} \cdot x)}{\left(\sum_{i=1}^K e^{u_i x + v_i} \right)^2} \\
&= \frac{e^{u_k x + v_k} \cdot x \cdot \left(\sum_{i=1}^K e^{u_i x + v_i} - e^{u_k x + v_k} \right)}{\left(\sum_{i=1}^K e^{u_i x + v_i} \right)^2} \\
&= w_k(x) \cdot x \cdot (1 - w_k(x))
\end{aligned} \tag{15}$$

2. Derivative of $w_k(x)$ with respect to v_k :

$$\begin{aligned}
\frac{\partial w_k(x)}{\partial v_k} &= \frac{\partial}{\partial v_k} \left(\frac{e^{u_k x + v_k}}{\sum_{i=1}^K e^{u_i x + v_i}} \right) \\
&= \frac{e^{u_k x + v_k} \cdot \left(\sum_{i=1}^K e^{u_i x + v_i} \right) - e^{u_k x + v_k} \cdot \frac{\partial}{\partial v_k} \left(\sum_{i=1}^K e^{u_i x + v_i} \right)}{\left(\sum_{i=1}^K e^{u_i x + v_i} \right)^2} \\
&= \frac{e^{u_k x + v_k} \cdot \left(\sum_{i=1}^K e^{u_i x + v_i} \right) - e^{u_k x + v_k} \cdot (e^{u_k x + v_k})}{\left(\sum_{i=1}^K e^{u_i x + v_i} \right)^2} \\
&= \frac{e^{u_k x + v_k} \cdot \left(\sum_{i=1}^K e^{u_i x + v_i} - e^{u_k x + v_k} \right)}{\left(\sum_{i=1}^K e^{u_i x + v_i} \right)^2} \\
&= w_k(x) \cdot (1 - w_k(x))
\end{aligned} \tag{16}$$

3. Derivative of $w_j(x)$ with respect to u_k (for $j \neq k$):

$$\begin{aligned}
\frac{\partial w_j(x)}{\partial u_k} &= \frac{\partial}{\partial u_k} \left(\frac{e^{u_j x + v_j}}{\sum_{i=1}^K e^{u_i x + v_i}} \right) \\
&= - \frac{e^{u_j x + v_j}}{\left(\sum_{i=1}^K e^{u_i x + v_i} \right)^2} \cdot e^{u_k x + v_k} \cdot x \\
&= -w_j(x) \cdot w_k(x) \cdot x
\end{aligned} \tag{17}$$

4. Derivative of $w_j(x)$ with respect to v_k (for $j \neq k$):

$$\begin{aligned}
 \frac{\partial w_j(x)}{\partial v_k} &= \frac{\partial}{\partial v_k} \left(\frac{e^{u_j x + v_j}}{\sum_{i=1}^K e^{u_i x + v_i}} \right) \\
 &= - \frac{e^{u_j x + v_j}}{\left(\sum_{i=1}^K e^{u_i x + v_i} \right)^2} \cdot e^{u_k x + v_k} \\
 &= -w_j(x) \cdot w_k(x)
 \end{aligned} \tag{18}$$

Using the results obtained from the derivatives above, we can calculate the following partial derivatives:

1. For u_k :

$$\begin{aligned}
 J_{n,k} &= \frac{\partial \hat{y}(x_n)}{\partial u_k} = \frac{\partial}{\partial u_k} \left(\sum_{k=1}^K w_k(x_n) \hat{y}_k(x_n) \right) \\
 &= \hat{y}_k(x_n) \frac{\partial w_k(x_n)}{\partial u_k} + \sum_{j \neq k} \hat{y}_j(x_n) \frac{\partial w_j(x_n)}{\partial u_k} \\
 &= \hat{y}_k(x_n) \cdot w_k(x_n) (1 - w_k(x_n)) \cdot x_n - \sum_{j \neq k} \hat{y}_j(x_n) \cdot w_j(x_n) w_k(x_n) \cdot x_n \\
 &= w_k(x_n) \cdot x_n \cdot \left(\hat{y}_k(x_n) (1 - w_k(x_n)) - \sum_{j \neq k} w_j(x_n) \cdot \hat{y}_j(x_n) \right)
 \end{aligned} \tag{19}$$

2. For v_k :

$$\begin{aligned}
 J_{n,(K-1)+k} &= \frac{\partial \hat{y}(x_n)}{\partial v_k} = \frac{\partial}{\partial v_k} \left(\sum_{k=1}^K w_k(x_n) \hat{y}_k(x_n) \right) \\
 &= \hat{y}_k(x_n) \frac{\partial w_k(x_n)}{\partial v_k} + \sum_{j \neq k} \hat{y}_j(x_n) \frac{\partial w_j(x_n)}{\partial v_k} \\
 &= \hat{y}_k(x_n) \cdot w_k(x_n) (1 - w_k(x_n)) - \sum_{j \neq k} \hat{y}_j(x_n) \cdot w_j(x_n) w_k(x_n) \\
 &= w_k(x_n) \cdot \left(\hat{y}_k(x_n) (1 - w_k(x_n)) - \sum_{j \neq k} w_j(x_n) \cdot \hat{y}_j(x_n) \right)
 \end{aligned} \tag{20}$$

3. For s_k :

$$\begin{aligned}
 J_{n,2(K-1)+k} &= \frac{\partial \hat{y}(x_n)}{\partial s_k} = \frac{\partial}{\partial s_k} (w_k(x_n) \cdot (s_k x_n + r_k)) \\
 &= w_k(x_n) \cdot \frac{\partial}{\partial s_k} (s_k x_n + r_k) \\
 &= w_k(x_n) \cdot x_n
 \end{aligned} \tag{21}$$

4. For r_k :

$$\begin{aligned}
 J_{n,(2(K-1)+K)+k} &= \frac{\partial \hat{y}(x_n)}{\partial r_k} = \frac{\partial}{\partial r_k} (w_k(x_n) \cdot (s_k x_n + r_k)) \\
 &= w_k(x_n) \cdot \frac{\partial}{\partial r_k} (s_k x_n + r_k) \\
 &= w_k(x_n)
 \end{aligned} \tag{22}$$

Finally, we can summarize the gradient as:

$$\nabla f = 2 \sum_{n=1}^N (\hat{y}(x_n) - y_n) J \tag{23}$$

10.1 Results for Optimization Variables

After implementing the LM method and running the optimization, the algorithm converged to the following values for the variables:

$$u = \begin{bmatrix} -13.077 \\ -80.105 \\ 88.202 \end{bmatrix} \quad v = \begin{bmatrix} -2.346 \\ -415.562 \\ -184.845 \end{bmatrix} \quad s = \begin{bmatrix} -1.202 \\ 3.265 \\ -4.000 \\ 12.126 \end{bmatrix} \quad r = \begin{bmatrix} -0.077 \\ 23.059 \\ 19.877 \\ -2.653 \end{bmatrix}$$

10.2 Plots

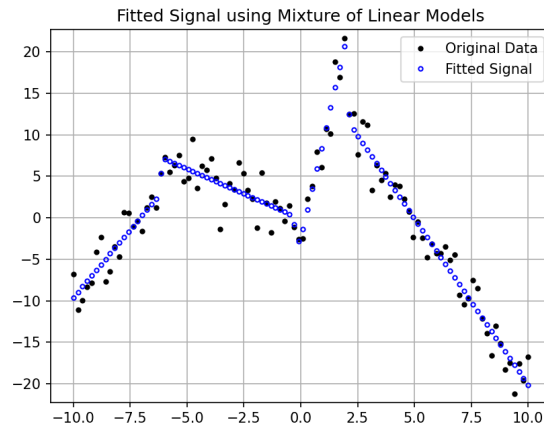


Figure 4: The fitted signal using the mixture of linear models. Each black dot represents a measurement, and the blue circles represent the output of the linear mixture.

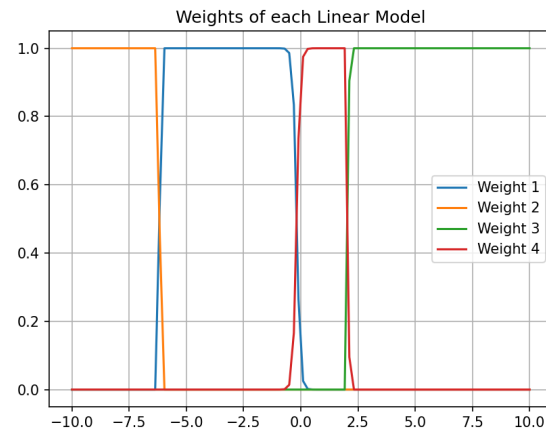


Figure 5: The weights $w_k(x)$ for each linear model. Each color represents a different weight function.




Figure 6: Values of the objective function across the LM iterations.




Figure 7: The norm of the gradient of the objective function across the LM iterations.

10.3 Comment on the Results

The Levenberg-Marquardt method converged effectively after 78 iterations, successfully bringing the gradient norm below the set threshold of $\epsilon = 10^{-4}$. As a result, the fitted signal aligns closely with the noisy measurements we started with. The weight functions behaved as anticipated, with each model activating in different regions of the signal, contributing to the overall fit. Throughout the optimization process, the objective function consistently decreased, demonstrating strong convergence behavior, and we observed a clear reduction in the gradient norm over time.

In the end, the final fitted parameters provided an accurate representation of the signal, effectively capturing the distinct linear segments and integrating them into a coherent overall model.

11 Code Developed

11.1 Python Code For Task 6

```

1 import cvxpy as cp
2 import numpy as np
3 import scipy.io
4 import matplotlib.pyplot as plt
5
6 # Open the data file .mat
7 data = scipy.io.loadmat('classifier_dataset.mat')
8
9 # Extract the variables from the data
10 traindataset = data['traindataset']
11 testdataset = data['testdataset']
12
13 trainlabels = data['trainlabels'].flatten() # Flatten to convert into 1D
14          array
15 testlabels = data['testlabels'].flatten() # Flatten to convert into 1D
16          array
17
18 # Save the number of rows in N (number of samples) and columns in D (number
19   of features)
20 N, D = traindataset.shape
21
22 # Regularization parameter
23 ro = 0.1
24
25 # Define the optimization variables
26 w0 = cp.Variable() # Bias term (scalar)
27 w = cp.Variable(D) # Weights vector (D-dimensional)
28
29 # Define the hinge loss function  $h(u) = \max(0, 1 - u)$ 
30 hinge_losses = cp.pos(1 - cp.multiply(trainlabels, traindataset @ w + w0))
31
32 # Define the objective function: hinge loss + regularization term
33 objective = cp.Minimize((1/N) * cp.sum(hinge_losses) + ro * cp.norm(w, 2)
34   **2)
35
36 # Define the problem and solve it
37 problem = cp.Problem(objective)
38 problem.solve()
39
40 # Retrieve the optimal parameters
41 w_optimal = w.value
42 w0_optimal = w0.value
43
44 # Function to evaluate the classifier error rate fD on a given dataset
45 def evaluate_error_rate(dataset, labels, w0_opt, w_opt):
46     predictions = np.sign(dataset @ w_opt + w0_opt) # Classify
47     misclassifications = np.sum(predictions != labels) # Count errors
48     error_rate = misclassifications / len(labels) # Calculate error rate
49     return error_rate

```

```
47 # Evaluate the classifier error rate on the training dataset
48 train_error_rate = evaluate_error_rate(traindataset, trainlabels, w0_optimal
    , w_optimal)
49 print(f"Training dataset error rate: {train_error_rate * 100:.2f}%")
50
51 # Evaluate the classifier error rate on the test dataset
52 test_error_rate = evaluate_error_rate(testdataset, testlabels, w0_optimal,
    w_optimal)
53 print(f"Test dataset error rate: {test_error_rate * 100:.2f}%")
```

Listing 1: Python Code to Solve Task 6

11.2 Python Code For Task 8

```

1 import cvxpy as cp
2 import numpy as np
3 import scipy.io
4 import matplotlib.pyplot as plt
5
6 # Open the data file .mat
7 data = scipy.io.loadmat('classifier_dataset.mat')
8
9 # Extract the variables from the data
10 traindataset = data['traindataset']
11 testdataset = data['testdataset']
12
13 trainlabels = data['trainlabels'].flatten() # Flatten to convert into 1D
14          array
15 testlabels = data['testlabels'].flatten() # Flatten to convert into 1D
16          array
17
18 # Save the number of rows in N (number of samples) and columns in D (number
19   of features)
20 N, D = traindataset.shape
21
22 # Regularization parameter
23 ro = 0.1
24
25 # Define the optimization variables
26 w0 = cp.Variable() # Bias term (scalar)
27 w = cp.Variable(D) # Weights vector (D-dimensional)
28
29 # Define the hinge loss function  $h(u) = \max(0, 1 - u)$ 
30 hinge_losses = cp.pos(1 - cp.multiply(trainlabels, traindataset @ w + w0))
31
32 # Define the objective function: hinge loss + regularization term
33 objective = cp.Minimize((1/N) * cp.sum(hinge_losses) + ro * cp.norm(w, 2)
34   **2)
35
36 # Define the problem and solve it
37 problem = cp.Problem(objective)
38 problem.solve()
39
40 # Retrieve the optimal parameters
41 w_optimal = w.value
42 w0_optimal = w0.value
43
44 # Function to evaluate the classifier error rate fD on a given dataset
45 def evaluate_error_rate(dataset, labels, w0_opt, w_opt):
46     predictions = np.sign(dataset @ w_opt + w0_opt) # Classify
47     misclassifications = np.sum(predictions != labels) # Count errors
48     error_rate = misclassifications / len(labels) # Calculate error rate
49     return error_rate

```

```

50 #####
51 #make the  $\tilde{x}$  vector (attacker vector)
52 #####
53 N2, D2 = testdataset.shape
54
55 x_attack_final=np.empty(shape=(N2, D2))
56
57 #For each sample, calculate  $\tilde{x}$ 
58 for k in range(0,N2):
59     #calculate  $y*w$ 
60     yw=testlabels[k]*w_optimal
61     i=0
62
63     #aply sign function to all elements in the vector yw
64     for num in yw:
65         if num>=0:
66             yw[i]= 1
67         else:
68             yw[i]= -1
69         i+=1
70
71     x=testdataset[k]
72     #calculate  $\tilde{x}$  for this sample
73     x_attack= x - P*yw
74     #add to matrix with all  $\tilde{x}$ 
75     x_attack_final[k]=x_attack
76
77 #####
78 #calculate error with attacked input
79 #####
80
81 test_error_rate = evaluate_error_rate(x_attack_final, testlabels, w0_optimal
82                                     , w_optimal)
83
84 print(f"Test dataset error rate with attack vector: {test_error_rate *
85       100:.2f}%")

```

Listing 2: Python Code to Solve Task 8

11.3 Python Code For Task 9

```

1
2 import cvxpy as cp
3 import numpy as np
4 import scipy.io
5 import matplotlib.pyplot as plt
6
7
8
9 # Open the data file .mat
10 data = scipy.io.loadmat('classifier_dataset.mat')
11
12 # Extract the variables from the data
13 traindataset = data['traindataset']
14 testdataset = data['testdataset']
15
16
17 trainlabels = data['trainlabels'].flatten() # Flatten to convert into 1D
18          array
19 testlabels = data['testlabels'].flatten()   # Flatten to convert into 1D
20          array
21
22 # Save the number of rows in N (number of samples) and columns in D (number
23   of features)
24 N, D = traindataset.shape
25
26 # Regularization parameter
27 ro = 0.1
28
29 # Define the optimization variables
30 w0 = cp.Variable() # Bias term (scalar)
31 w = cp.Variable(D) # Weights vector (D-dimensional)
32
33 # Define the hinge loss function  $h(u) = \max(0, 1 - u)$ 
34 hinge_losses = cp.pos(1 - cp.multiply(trainlabels, traindataset @ w + w0))
35
36 # Define the objective function: hinge loss + regularization term
37 objective = cp.Minimize((1/N) * cp.sum(hinge_losses) + ro * cp.norm(w, 2)
38                          **2)
39
40 # Define the problem and solve it
41 problem = cp.Problem(objective)
42 problem.solve()
43
44 # Retrieve the optimal parameters
45 w_optimal = w.value
46 w0_optimal = w0.value
47
48 # Function to evaluate the classifier error rate fD on a given dataset
49 def evaluate_error_rate(dataset, labels, w0_opt, w_opt):

```

```

50     predictions = np.sign(dataset @ w_opt + w0_opt) # Classify
51     misclassifications = np.sum(predictions != labels) # Count errors
52     error_rate = misclassifications / len(labels) # Calculate error rate
53     return error_rate
54
55
56 P=0.18
57
58 #####
59 #make the  $x^{\sim}$  vector (attacker vector)
60 #####
61 N2, D2 = testdataset.shape
62
63 x_attack_final=np.empty(shape=(N2, D2))
64
65
66
67 #For each sample, calculate  $x^{\sim}$ 
68 for k in range(0,N2):
69     #calculate  $y*w$ 
70
71     yw=testlabels[k]*w_optimal
72     i=0
73
74     #aply sign function to all elements in the vector yw
75     for num in yw:
76         if num>=0:
77             yw[i]= 1
78         else:
79             yw[i]= -1
80         i+=1
81
82     x=testdataset[k]
83     #calculate  $x^{\sim}$  for this sample
84     x_attack= x - P*yw
85     #add to matrix with all  $x^{\sim}$ 
86     x_attack_final[k]=x_attack
87
88
89
90 w0 = cp.Variable()
91 w = cp.Variable(D)
92
93 ro = 0.1
94
95
96 #Vector manipulation
97 train_labels_col = cp.reshape(trainlabels, (400,1))
98 w_T = cp.reshape(w, (1,784))
99
100 #Compute yw
101 product_matrix = train_labels_col @ w_T
102
103 #Equivalent to calculating the l1-norm for every row
104 l1_norm = cp.sum(cp.abs(product_matrix), axis=1)

```



```
105
106 #Define the hinge loss function
107 hinge_losses_9 = cp.pos(1 - (cp.multiply(trainlabels, traindataset @ w + w0)
    - P * l1_norm))
108
109 #Problem definition
110 objective_9 = cp.Minimize((1/N) * cp.sum(hinge_losses_9) + ro * cp.norm(w,
    2)**2)
111 problem_9 = cp.Problem(objective_9)
112 problem_9.solve()
113
114 w_optimal = w.value
115 w0_optimal = w0.value
116
117
118 #Result Evaluation
119 train_error_rate = evaluate_error_rate(traindataset, trainlabels, w0_optimal
    , w_optimal)
120 print(f"Training dataset error rate: {train_error_rate * 100:.2f}%")
121
122 test_error_rate = evaluate_error_rate(testdataset, testlabels, w0_optimal,
    w_optimal)
123 print(f"Test dataset error rate: {test_error_rate * 100:.2f}%")
124
125 attack_test_error_rate = evaluate_error_rate(x_attack_final, testlabels,
    w0_optimal, w_optimal)
126 print(f"Test dataset error rate with attack vector: {attack_test_error_rate
    * 100:.2f}%")
```

Listing 3: Python Code to Solve Task 9

11.4 Python Code For Task 10

```

1 import numpy as np
2 import scipy.io as sio
3 import matplotlib.pyplot as plt
4
5
6 # Load the .mat file
7 def load_data(file_path):
8     data = sio.loadmat(file_path)
9     X = data['x'] # Input data (X values)
10    Y = data['y'] # Target data (Y values)
11    U = data['u'] # ...
12    V = data['v'] # ...
13    S = data['s'] # ...
14    R = data['r'] # ...
15    return X.flatten(), Y.flatten(), U.flatten(), V.flatten(), S.flatten(),
16    R.flatten()
17
18
19 # Mixture of Linear Models
20 def mixture_model(x, u, v, s, r, K, N):
21     y_pred = np.zeros(N)
22     W = np.zeros((N, K))
23
24     alphas = np.zeros((N, K))
25     for k in range(K-1):
26         alphas[:, k] = u[k] * x + v[k]
27
28     # Centering to prevent overflow
29     max_alpha = np.max(alphas, axis=1, keepdims=True)
30     exp_alphas = np.exp(alphas - max_alpha)
31
32     # Normalizing weights
33     for k in range(K):
34         W[:, k] = exp_alphas[:, k] / np.sum(exp_alphas, axis=1)
35
36     # Combine predictions from all K models
37     for k in range(K):
38         y_pred += W[:, k] * (s[k] * x + r[k])
39
40     return y_pred, W
41
42
43
44 # Compute partial derivatives (Jacobian elements)
45 def compute_jacobian_partial(x, u, v, s, r, W, k, K, N, param):
46     if param == 's':
47         return W[:, k] * x
48     elif param == 'r':
49         return W[:, k]
50     elif param == 'u':
51         J_partial_u = np.zeros(N)
52         for i in range(N):

```

```

53         k_term = W[i, k] * (1 - W[i, k]) * x[i]*(s[k] * x[i] + r[k])#
54 derivative of Wk*Yk
55
56         j_term = W[i, k]*x[i]*sum(W[i, j] * (s[j] * x[i] + r[j]) for j
57 in range(K) if j != k) # sum of derivative of Wj*Yj
58
59         J_partial_u[i] = k_term - j_term
60
61     return J_partial_u
62
63     elif param == 'v':
64         J_partial_v = np.zeros(N)
65         for i in range(N):
66
67             W_term = W[i, k] * (1 - W[i, k])*(s[k] * x[i] + r[k])#
68 derivative of Wj*Yj
69
70             sum_term = W[i, k]*sum(W[i, j] * (s[j] * x[i] + r[j]) for j in
71 range(K) if j != k) #sum of derivative Wj*Yj
72
73             J_partial_v[i] = W_term - sum_term
74
75         return J_partial_v
76     else:
77         raise ValueError(f"Invalid parameter: {param}")
78
79 # Compute the gradient and Jacobian
80 def compute_Jacobian(x, u, v, s, r, W, K, N):
81     J = np.zeros((N, 4 * K - 2)) # Jacobian
82
83     # Compute partial derivatives for uk and vk
84     for k in range(K - 1):
85         J[:, k] = compute_jacobian_partial(x, u, v, s, r, W, k, K, N, 'u')
86     # Jacobian for u_k
87     J[:, (K - 1) + k] = compute_jacobian_partial(x, u, v, s, r, W, k, K,
88 N, 'v') # Jacobian for v_k
89
90     # Compute partial derivatives for sk and rk
91     for k in range(K):
92         J[:, 2*(K - 1)+ k] = compute_jacobian_partial(x, u, v, s, r, W, k, K
93 , N, 's') # Jacobian for s_k
94         J[:, 2*(K - 1)+ K + k] = compute_jacobian_partial(x, u, v, s, r, W,
95 k, K, N, 'r') # Jacobian for r_k
96
97     return J
98
99 # Levenberg-Marquardt optimization
100 def levenberg_marquardt(x, y, u, v, s, r, K, N, max_iter=5000, epsilon=1e-4,
101 lambda_init=1.0):
102     lambda_ = lambda_init # Damping parameter
103     residuals_list = [] # To store sum of squared residuals for plotting
104     grad_obj_func_list = [] # To store gradient of objective function for

```

plotting

```

99
100     for iter in range(max_iter):
101
102         # Compute model prediction and weights
103         y_pred, W = mixture_model(x, u, v, s, r, K, N)
104         # Combine the u, v, s, r vectors into a single column vector
105         param_vector = np.concatenate([u, v, s, r]).reshape(-1, 1)
106
107         # Compute the residual and objective function
108         residual = y_pred - y
109         obj_func = np.sum(residual ** 2)
110         residuals_list.append(obj_func)
111         # Compute the full Jacobian matrix (gradients)
112         J = compute_Jacobian(x, u, v, s, r, W, K, N)
113
114         # Construct the A matrix (Jacobian and regularization term)
115         sqrt_lambda = np.sqrt(lambda_)
116
117         # Create a (4*K-2)x(4*K-2) identity matrix
118         identity_matrix = np.identity(4*K-2)
119
120         # Multiply the identity matrix by sqrt_lambda to get the final
121         sqrt_lambda matrix
122         sqrt_lambda_matrix = sqrt_lambda * identity_matrix
123
124         # Combine the jacobian with the square root of the scalar
125         A = np.vstack([J, sqrt_lambda_matrix])
126
127         J_param = J @ param_vector # Multiply jacobian by parameters vector
128         b_top = J_param - residual.reshape(-1, 1) # Subtract residual from
129         each row
130
131         # Bottom part: sqrt_lambda * param_vector
132         b_bottom = sqrt_lambda * identity_matrix @ param_vector # Shape
133         will be 14 x 1
134
135         # Combine the top and bottom parts
136         b = np.vstack([b_top, b_bottom])
137
138         # Solve the least-squares problem for min
139         minimized_parameters, _, _, _ = np.linalg.lstsq(A, b, rcond=None)
140
141         # Compute the candidate parameters with min
142         u_new, v_new, s_new, r_new = update_parameters(u, v, s, r,
143         minimized_parameters, K)
144         # Compute new prediction and objective function with updated
145         parameters
146         y_pred_new, W_new = mixture_model(x, u_new, v_new, s_new, r_new, K,
147         N)
148         obj_func_new = np.sum((y_pred_new - y) ** 2)
149
150         gradient = 2 * J.T @ residual

```

```

147     # Check stopping criterion
148     grad_obj_func_list.append(np.linalg.norm(gradient))
149     if np.abs(np.linalg.norm(gradient)) < epsilon:
150         print(f"Converged at iteration {iter}")
151         break
152     # Check if the step is valid
153
154     if obj_func_new < obj_func: # Valid step
155         u, v, s, r = u_new, v_new, s_new, r_new
156         lambda_ *= 0.7 # Decrease lambda
157
158     else: # Null step
159         lambda_ *= 2.0 # Increase lambda
160
161
162     # Plot the results
163
164     plot_results(x, y, u, v, s, r, y_pred, W, residuals_list,
165 grad_obj_func_list)
166     return u, v, s, r
167
168 # Update the parameters (u, v, s, r) with the delta step from LM
169 def update_parameters(u, v, s, r, minimized_parameters, K):
170
171     # Ensure the deltas are treated as 1D vectors instead of 2D
172     u_new = minimized_parameters[:K - 1].flatten() # Flatten in case it's
higher dimensional
173     v_new = minimized_parameters[K - 1:2 * K - 2].flatten()
174     s_new = minimized_parameters[2 * K - 2:3 * K - 2].flatten()
175     r_new = minimized_parameters[3 * K - 2:].flatten()
176
177     # Update the parameters with the corresponding deltas
178
179
180     return u_new, v_new, s_new, r_new
181
182 # Plot the fitted signal and the weights
183 def plot_results(x, y, u, v, s, r, y_pred, W, residuals_list,
grad_obj_func_list):
184
185     plt.figure()
186     plt.plot(x, y, 'ko', markersize=3, label='Original Data')
187     plt.plot(x, y_pred, 'o', markerfacecolor='none', markeredgecolor='blue',
markersize=3, label='Fitted Signal')
188     plt.ylim(min(y) - 1, max(y) + 1)
189     plt.xlim(min(x) - 1, max(x) + 1)
190     plt.title('Fitted Signal using Mixture of Linear Models')
191     plt.legend()
192     plt.grid(True)
193     plt.show()
194
195     plt.figure()
196     for k in range(W.shape[1]):
197         plt.plot(x, W[:, k], label=f'Weight {k+1}')

```

```

198 plt.title('Weights of each Linear Model')
199 plt.legend()
200 plt.grid(True)
201 plt.show()
202
203 plt.figure()
204 plt.plot(range(len(residuals_list)), residuals_list, 'b-', label='
Objective funtion')
205 plt.xlabel('Iterations')
206 plt.ylabel('Objective funtion')
207 plt.title('Objective funtion Over Iterations')
208 plt.legend()
209 plt.grid(True)
210 plt.yscale('log')
211 plt.show()
212
213 plt.figure()
214 plt.plot(range(len(grad_obj_func_list)), grad_obj_func_list, 'b-', label
='Gradient of objective funtion')
215 plt.xlabel('Iterations')
216 plt.ylabel('Gradient of objective funtion')
217 plt.title('Gradient of objective funtion Over Iterations')
218 plt.legend()
219 plt.ylim(0.00001, 100000)
220 plt.grid(True)
221 plt.yscale('log')
222 plt.show()
223
224
225
226
227 # Example usage
228 file_path = 'lm_dataset_task.mat' # Replace with actual path
229 # Initial values for u, v, s, r
230 x, y, u, v, s, r = load_data(file_path)
231
232
233 K = len(s) # Number of models
234 N = len(x) # Number of data points
235
236 # Run the LM optimization
237 u_opt, v_opt, s_opt, r_opt = levenberg_marquardt(x, y, u, v, s, r, K, N)
238 #print results
239 print(u_opt)
240 print(v_opt)
241 print(s_opt)
242 print(r_opt)

```

Listing 4: Python Code to Solve Task 10